

# GNN Implementation: PyG & GraphGym

**Jiaxuan You**  
**Assistant Professor at UIUC CDS**



**CS598: Deep Learning with Graphs, 2024 Fall**  
**<https://ulab-uiuc.github.io/CS598/>**

# Logistics: Proposal Task Due

- The deadline for submission is **Oct 6 (Sunday), 11:59 PM, CT.**
- Please submit a **PDF** file to Canvas for each group. Each group only needs to submit once. Only the last submission from each group will be graded.
- You are required to use the ICLR 2025 template. **Abstracts are not required.** Please contact the TA if you have trouble using LaTeX.
- Submissions are suggested to be between 1.5 to 2 pages in length, with a minimum of 1 page, excluding references.

# Logistics: Project Idea Discussion

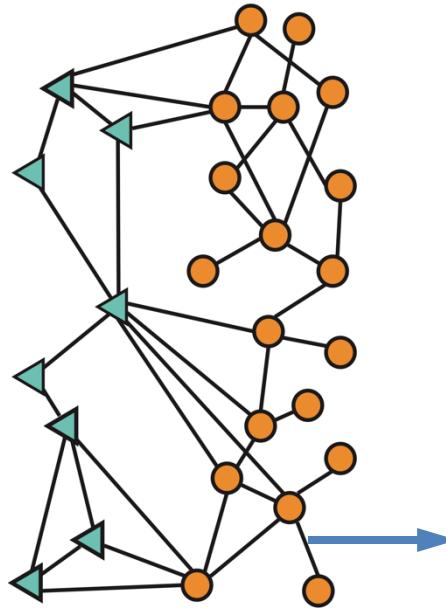
- **Wednesday Schedule**

- Each group is required to present your proposal (**around 3 minutes**) and answer questions from instructor and classmates (**around 1 minute**).
- Slides are **NOT** required but are welcomed.
- Share your proposal in Slack and provide your feedback to others' proposal after class.
  - As part of grading for “Ideate and discussion”, peer discussion on Slack will count towards **2%** of your final grade. Please provide feedback to **at least 2 groups**.
- **Attendance to each discussion session will count towards 1% of your final grade.**

# GNN Implementation

# GNN Design Space & GraphGym

# How to Design GNNs for a New Task?



A novel GNN application



J.P.Morgan

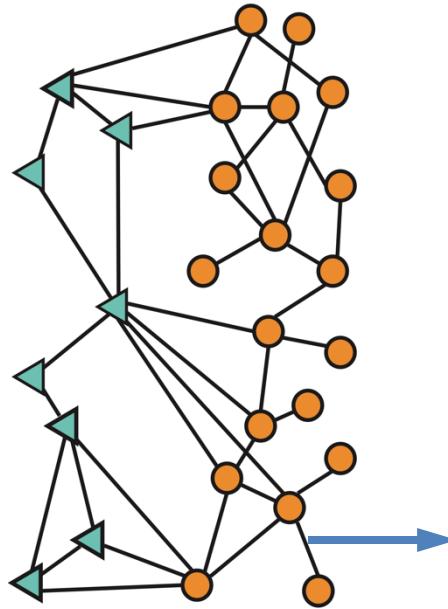
**Edge-level prediction**

*Is this transaction fraudulent?*

	ROC AUC
Non-graph baseline	0.8
GNN (out-of-the-box)	0.81

Graph ML seems to have no benefits

# How to Design GNNs for a New Task?



A novel GNN application



J.P.Morgan

Edge-level prediction

*Is this transaction fraudulent?*

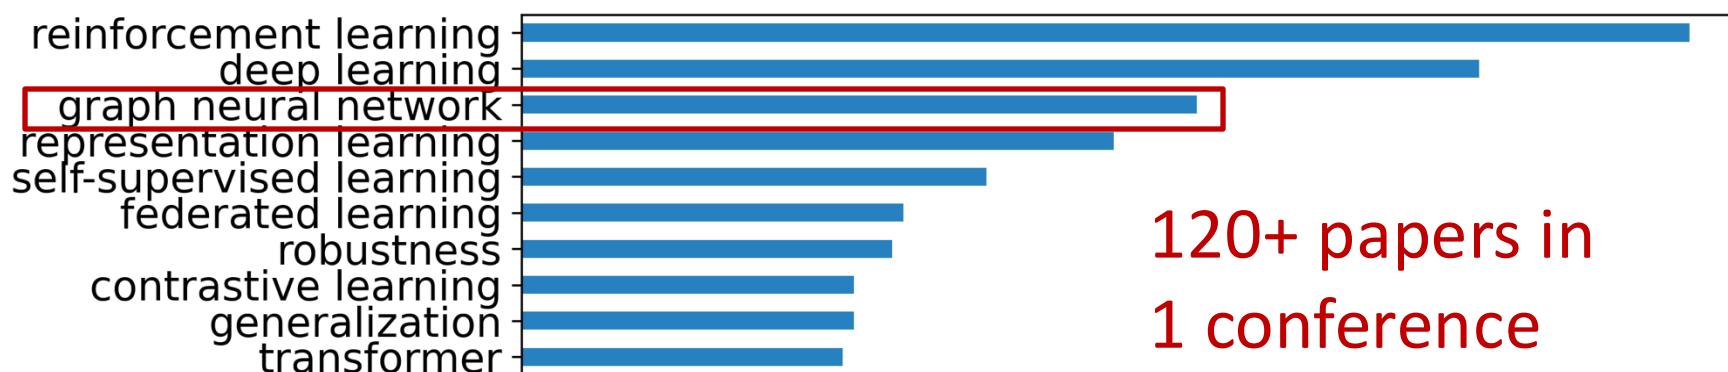
	ROC AUC
Non-graph baseline	0.8
GNN (out-of-the-box)	0.81
GNN (careful design)	0.94

- Question: **Principled way** to design a good GNN for a new task?
  - Promote the research & application of ML on graphs

# Designing GNNs is Challenging

- **Numerous GNN models** have been proposed – which one to use?

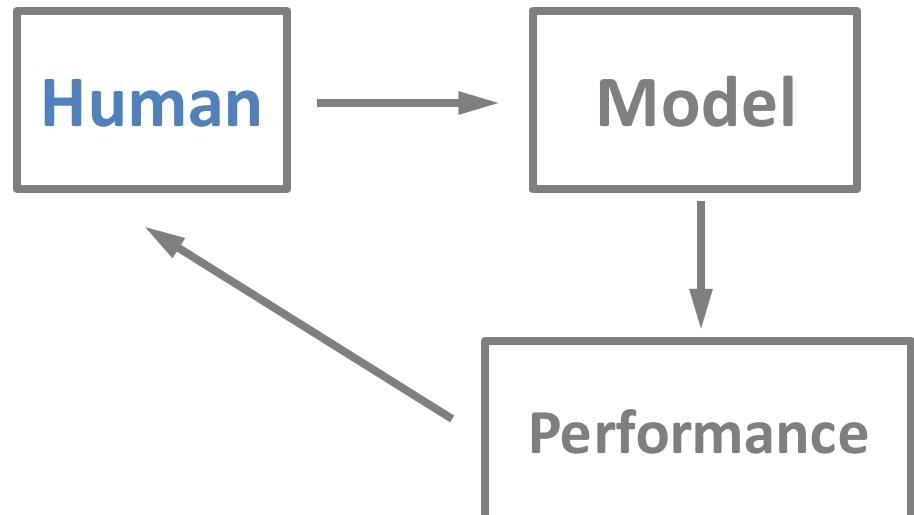
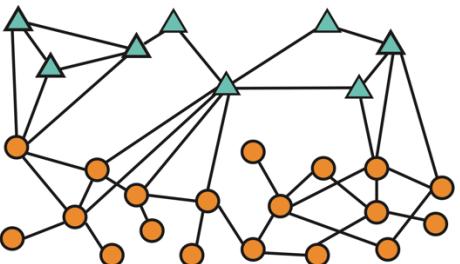
ICLR 2022 most appeared keywords



- **Issue: no principled approach** for finding the best GNN designs

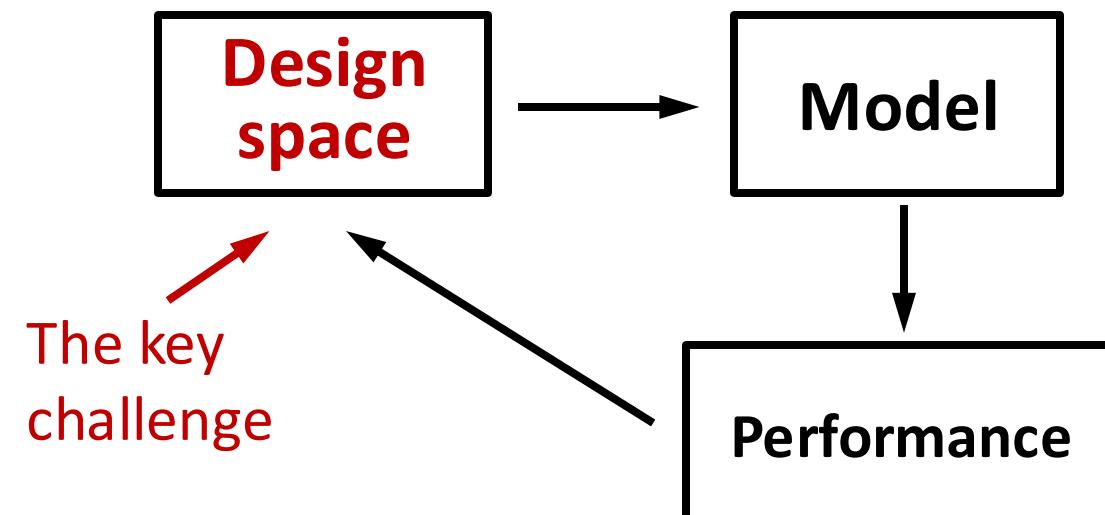
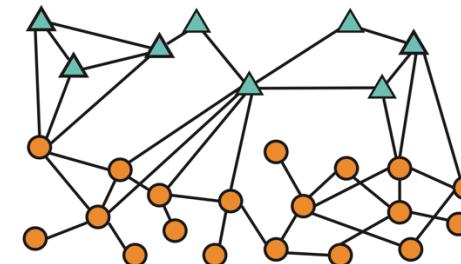
# Solution: AutoML on graphs

Graph  
ML task



Previous works

Graph  
ML task



My approach

# Background: Terminology

- **Design:** a concrete model instantiation
  - E.g., a 4-layer GraphSAGE
- **Design dimensions** characterize a design
  - E.g., the number of layers  $L \in \{2, 4, 6, 8\}$
- **Design choice** is the actual selected value in the design dimension
  - E.g., the number of layers  $L = 2$
- **Design space** consists of a Cartesian product of design dimensions
- **Task:** A specific task of interest
  - E.g., node classification on Cora, graph classification on ENZYMES
- **Task space** consists of all the tasks we care about

# Key Questions for GNN Design

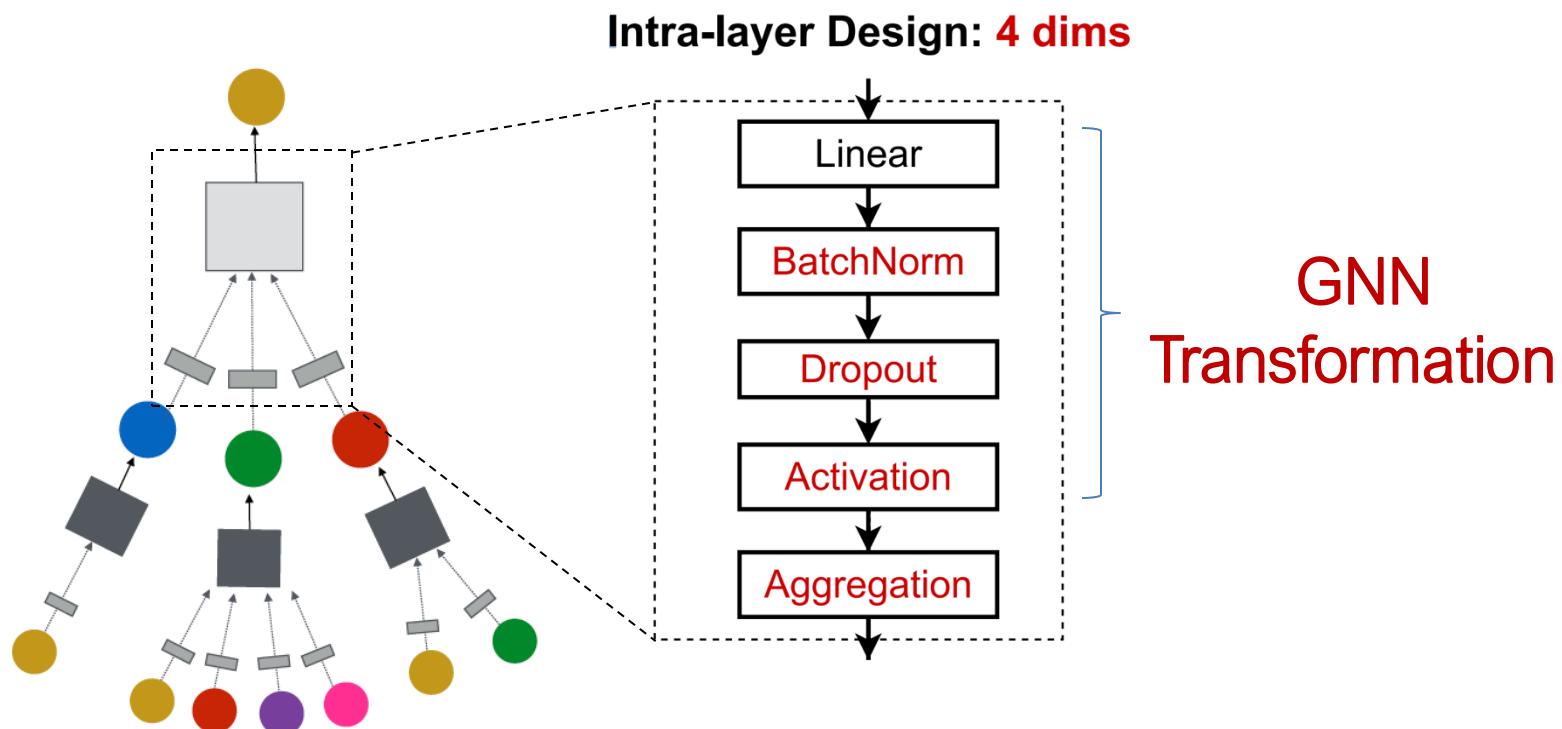
- **GNN architectural design:**
  - How to find a good GNN design for a specific GNN task?
- **Important but challenging:**
  - Domain experts want to use SOTA GNN on their specific tasks, however...
    - There are tons of possible GNN architectures
      - GCN, GraphSAGE, GAT, GIN, ...
    - **Issue:** Best design in one task can perform badly for another task
    - Redo hyperparameter grid search for each new task is NOT feasible
- **Key contribution of GraphGym paper:**
  - The first systematic study for the ***GNN design space and task space***
  - **GraphGym**, a powerful platform for exploring different GNN designs and tasks

# Recap: GNN Design Space

## Intra-layer Design:

**GNN Layer = Transformation + Aggregation**

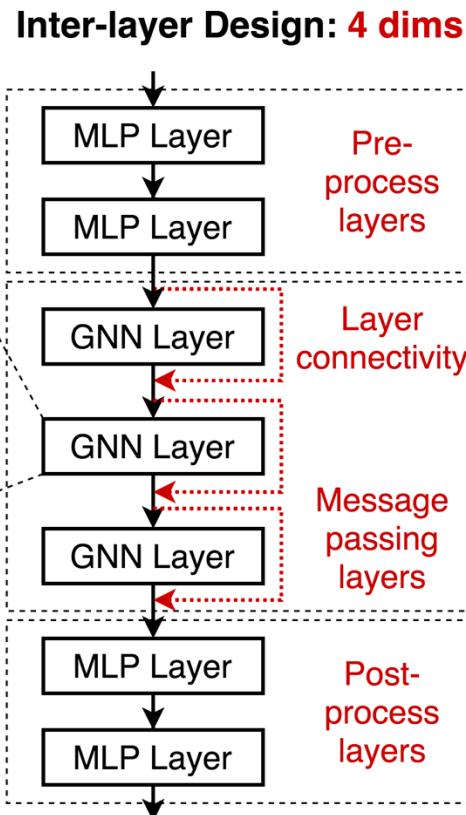
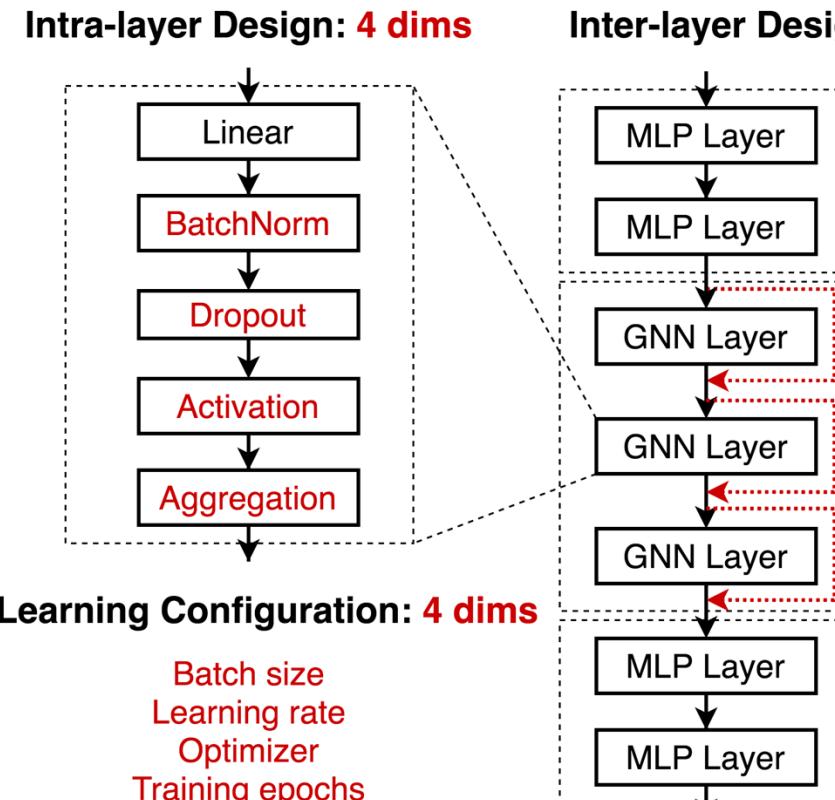
- We propose a general instantiation under this perspective



# Recap: GNN Design Space

## Inter-layer Design

- We explore different ways of organizing GNN layers



### Pre-process layers:

Important when expressive node feature encoder is needed

E.g., when nodes are images/text

### Skip connections:

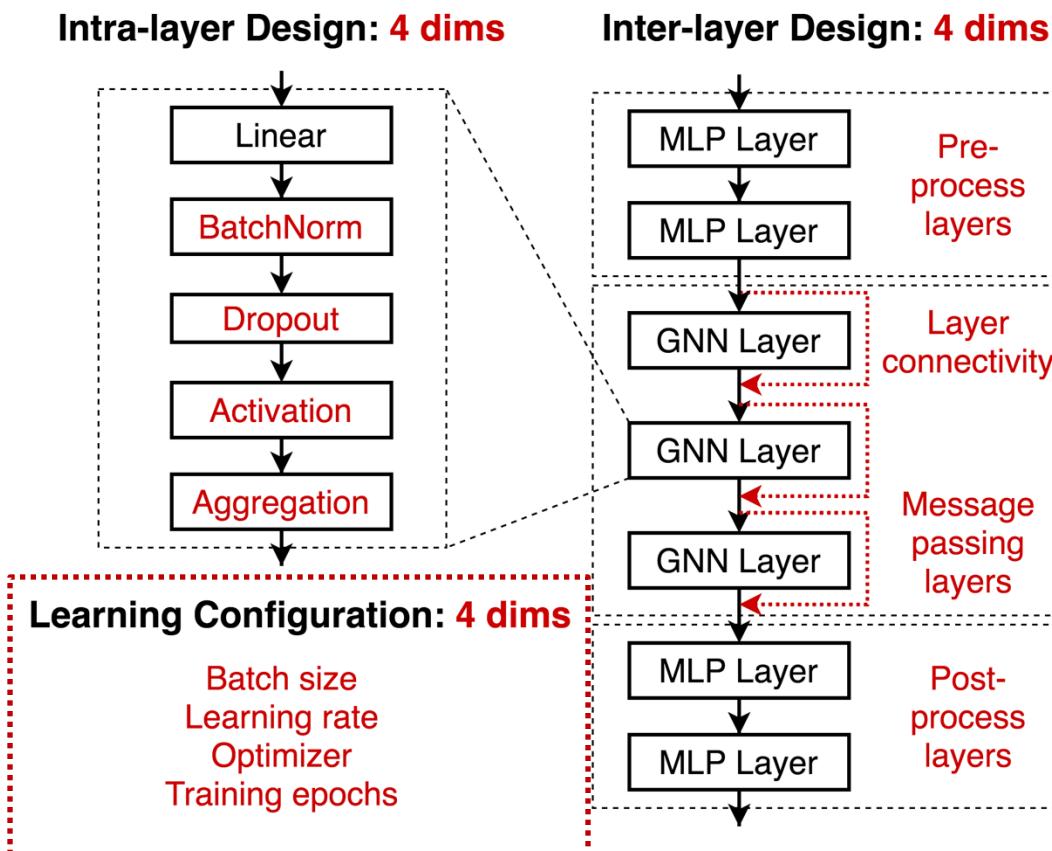
Improve deep GNN's performance

### Post-process layers:

Important when reasoning or transformation over node embeddings are needed

E.g., graph classification, knowledge graphs

# Recap: GNN Design Space



- ## Learning configurations
- Often neglected in current literature
  - But we found they have high impact on performance

# GNN Design Space

## ■ Overall: A GNN design space

### ■ Intra-layer design

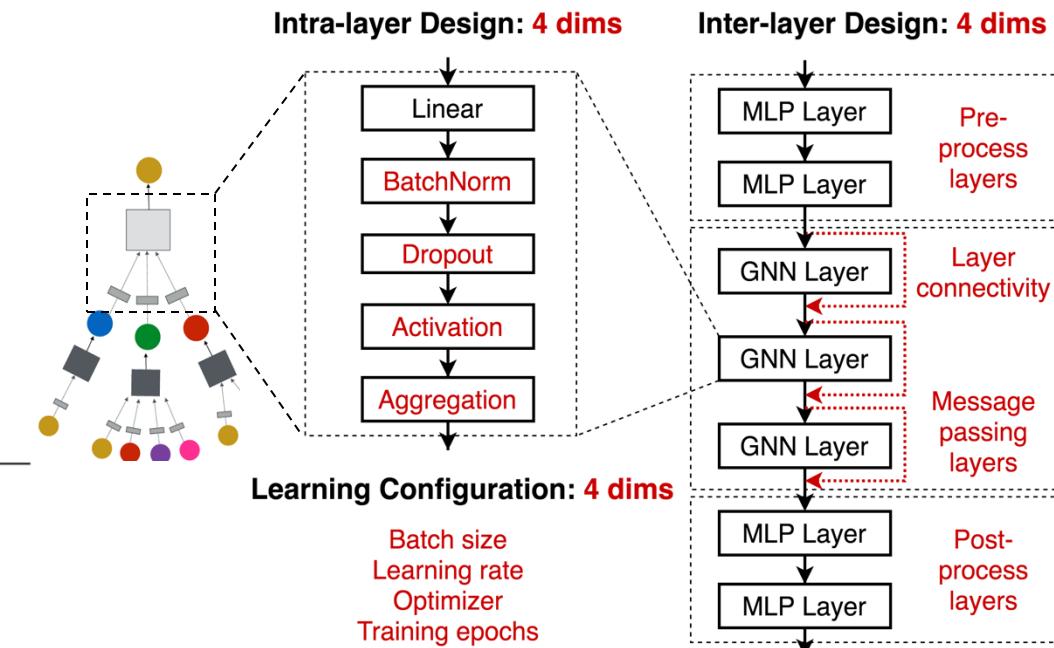
Batch Normalization	Dropout	Activation	Aggregation
True, False	False, 0.3, 0.6	RELU, PRELU, SWISH	MEAN, MAX, SUM
Layer connectivity	Pre-process layers	Message passing layers	Post-precess layer
STACK, SKIP-SUM, SKIP-CAT	1, 2, 3	2, 4, 6, 8	1, 2, 3
Batch size	Learning rate	Optimizer	Training epochs
16, 32, 64	0.1, 0.01, 0.001	SGD, ADAM	100, 200, 400

### ■ Learning configuration

### ■ In total: 315K possible designs

## ■ Purpose:

- We don't want to (and we cannot) cover all the possible designs
- A mindset transition: We want to demonstrate that **studying a design space is more effective than studying individual GNN designs**



# A General GNN Task Space

- **Categorizing GNN tasks**
  - **Common practice:** node / edge / graph level task
  - Reasonable but not precise
    - **Node prediction:** predict **clustering coefficient** vs. predict a **node's subject area in a citation networks** – **completely different task**
  - But creating a precise taxonomy of GNN tasks is very hard!
    - **Subjective; Novel GNN tasks** can always emerge
- **Our innovation: a quantitative task similarity metric**
  - **Purpose:** **understand GNN tasks, transfer the best GNN models across tasks**

# A General GNN Task Space

- **Innovation:** quantitative **task similarity metric**
  - 1) Select “**anchor**” models ( $M_1, \dots, M_5$ )
  - 2) Characterize a task by **ranking** the performance of anchor models
  - 3) Tasks with **similar rankings** are considered as similar

Task Similarity Metric

	Anchor Model Performance ranking					Similarity to Task A
Task A	$M_1$	$M_2$	$M_3$	$M_4$	$M_5$	1.0
Task B	$M_1$	$M_3$	$M_2$	$M_4$	$M_5$	0.8
Task C	$M_5$	$M_1$	$M_4$	$M_3$	$M_2$	-0.4

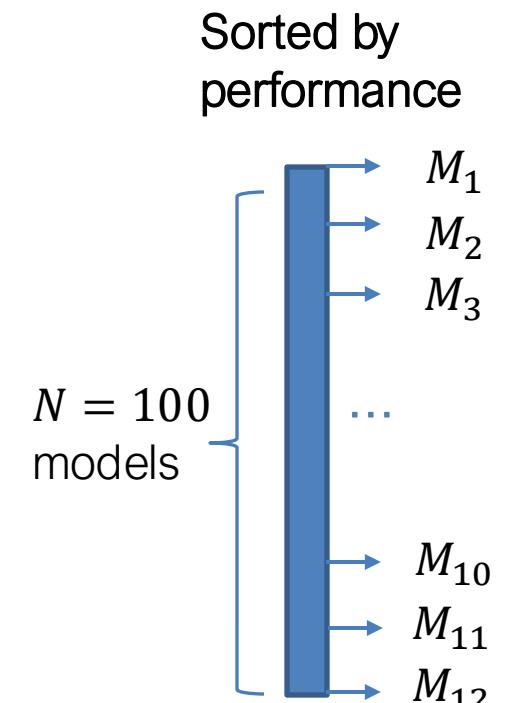
Task A is similar to Task B  
Task A is not similar to Task C

- How do we select the anchor models?

# A General GNN Task Space

## ■ Selecting the anchor models

- 1) Select a small dataset
  - E.g., node classification on Cora
- 2) Randomly sample  $N$  models from our design space, run on the dataset
  - E.g., we sample 100 models
- 3) Sort these models based on their performance:  
**evenly select  $M$  models as the anchor models**, whose performance range from the worst to the best
  - E.g., we sample 12 models in our experiments
- **Goal: Cover a wide spectrum of models:** A bad model in one task could be great for another task



# A General GNN Task Space

- We collect 32 tasks: node / graph classification

Task name
node-AmzonComputers-N/A-N/A
node-AmzonPhoto-N/A-N/A
node-CiteSeer-N/A-N/A
node-CoauthorCS-N/A-N/A
node-CoauthorPhysics-N/A-N/A
node-Cora-N/A-N/A
node-scalefree-clustering-pagerank
node-scalefree-const-clustering
node-scalefree-const-pagerank
node-scalefree-onehot-clustering
node-scalefree-onehot-pagerank
node-scalefree-pagerank-clustering
node-smallworld-clustering-pagerank
node-smallworld-const-clustering
node-smallworld-const-pagerank
node-smallworld-onehot-clustering
node-smallworld-onehot-pagerank
node-smallworld-pagerank-clustering
graph-PROTEINS-N/A-N/A
graph-BZR-N/A-N/A
graph-COX2-N/A-N/A
graph-DD-N/A-N/A
graph-ENYMES-N/A-N/A
graph-IMDB-N/A-N/A
graph-scalefree-clustering-path
graph-scalefree-const-path
graph-scalefree-onehot-path
graph-scalefree-pagerank-path
graph-smallworld-clustering-path
graph-smallworld-const-path
graph-smallworld-onehot-path
graph-smallworld-pagerank-path
graph-ogbg-molhiv-N/A-N/A

6 Real-world node classification tasks

12 Synthetic node classification tasks

Predict node properties:

- Clustering coefficient
- PageRank

6 Real-world graph classification tasks

8 Synthetic graph classification tasks

Predict graph properties:

- Average path length

# Evaluating GNN Designs

- **Evaluating a design dimension:**
  - “Is BatchNorm generally useful for GNNs?”
- **The common practice:**
  - (1) Pick one model (e.g., a 5-layer 64-dim GCN)
  - (2) Compare two models, with BN = True / False
- **Our approach:**
  - Note that **we have defined** 315K (models) \* 32 (tasks)  $\approx$  **10M model-task combinations**
  - (1) **Sample from 10M possible model-task combinations**
  - (2) **Rank the models** with BN = True / False
- How do we make it **scalable & convincing?**

# Evaluating GNN Designs

- Evaluating a design dimension: controlled random search
  - a) Sample random model-task configurations, perturb BatchNorm = [True, False]
  - Here we control the computational budget for all the models

(a) Controlled Random Search

GNN Design Space					GNN Task Space	
BatchNorm	Activation	...	Message layers	Layer Connectivity	Task level	dataset
True	relu	...	8	skip_sum	node	CiteSeer
False	relu	...	8	skip_sum	node	CiteSeer
True	relu	...	2	skip_cat	graph	BZR
False	relu	...	2	skip_cat	graph	BZR
...						
True	prelu	...	4	stack	graph	scale free
False	prelu	...	4	stack	graph	scale free

# Evaluating GNN Designs

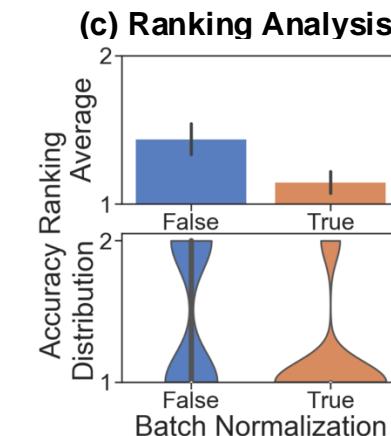
- **b) Rank** BatchNorm = [True, False] by their performance (lower ranking is better)
- **c) Plot Average / Distribution of the ranking** of BatchNorm = [True, False]

(b) Rank Design Choices by Performance

GNN Design Space	
BatchNorm	
True	0.75
False	0.54
True	0.88
False	0.88
True	0.89
False	0.36

Experimental Results

Val. Accuracy	Design Choice Ranking
0.75	1
0.54	2
0.88	1 (a tie)
0.88	1 (a tie)
0.89	1
0.36	2



- **Summary:** Convincingly evaluate any new design dimension, e.g., evaluate a new GNN layer we propose

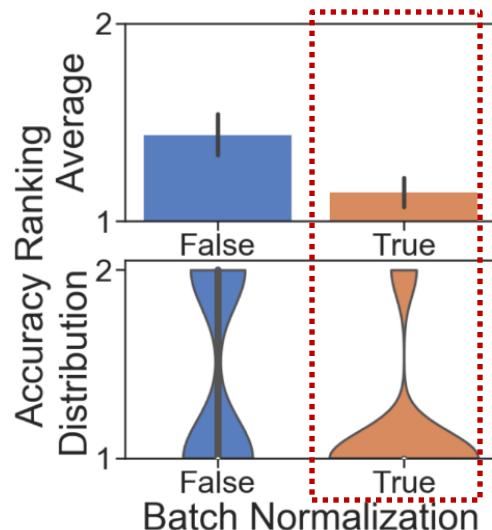
# Results 1: A Guideline for GNN Design

- Certain design choices exhibit **clear advantages**

- Intra-layer designs:

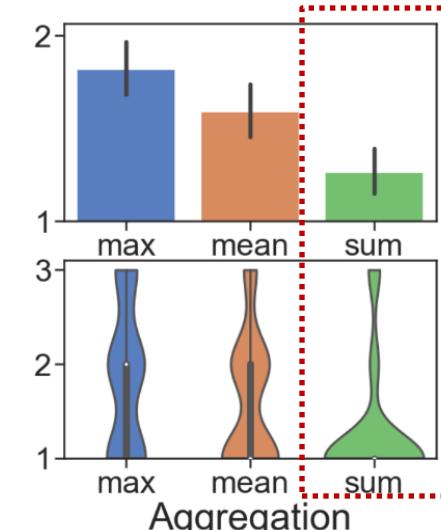
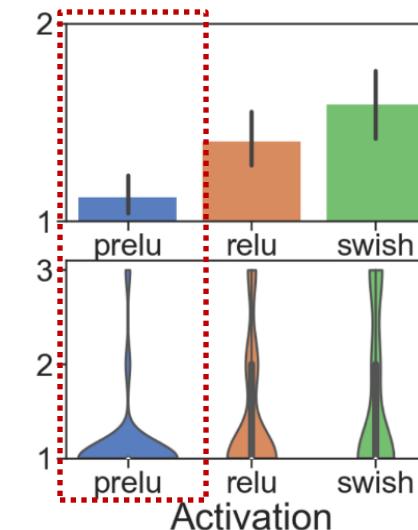
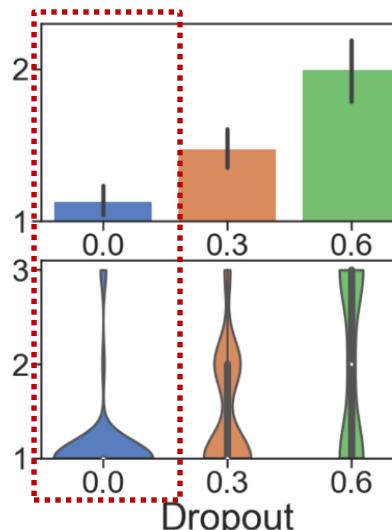
Explanation:

GNNs are hard to optimize



Explanation:

This is our new finding!



Explanation:

GNNs experience underfitting more often

Explanation:

Sum is the most expressive aggregator

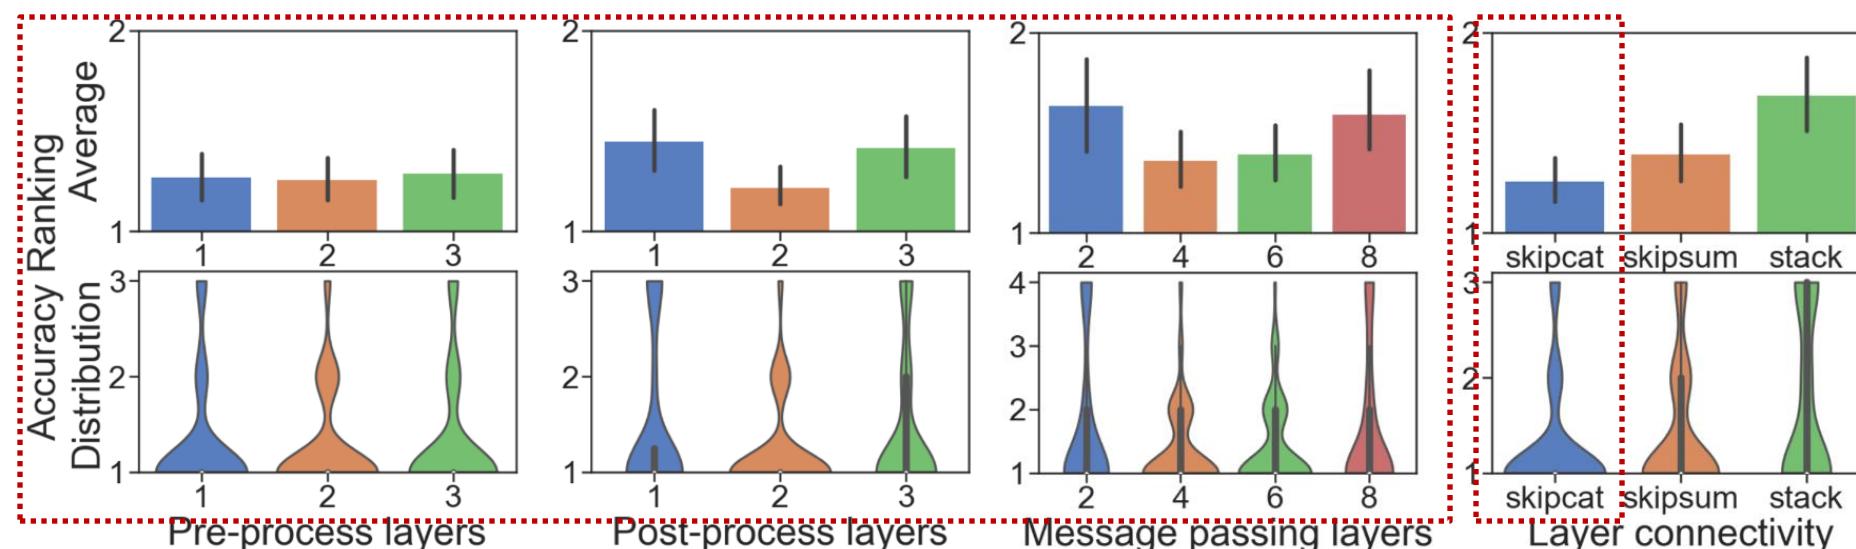
# Results 1: A Guideline for GNN Design

- Certain design choices exhibit **clear advantages**

- **Inter-layer designs**

Optimal number of layers is hard to decide

Highly dependent on the task



**Explanation:**  
Skip connection enable  
hierarchical node  
representation

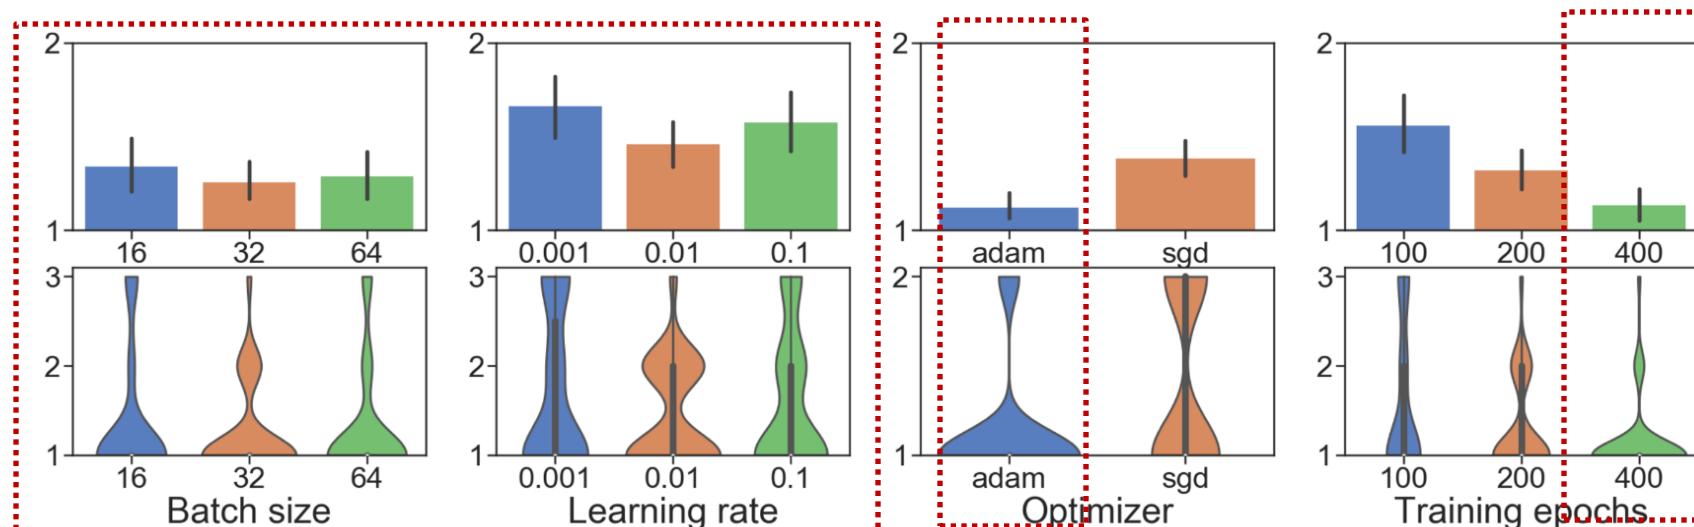
# Results 1: A Guideline for GNN Design

- Certain design choices exhibit **clear advantages**

- Learning configurations**

Optimal batch size and learning rate is hard to decide

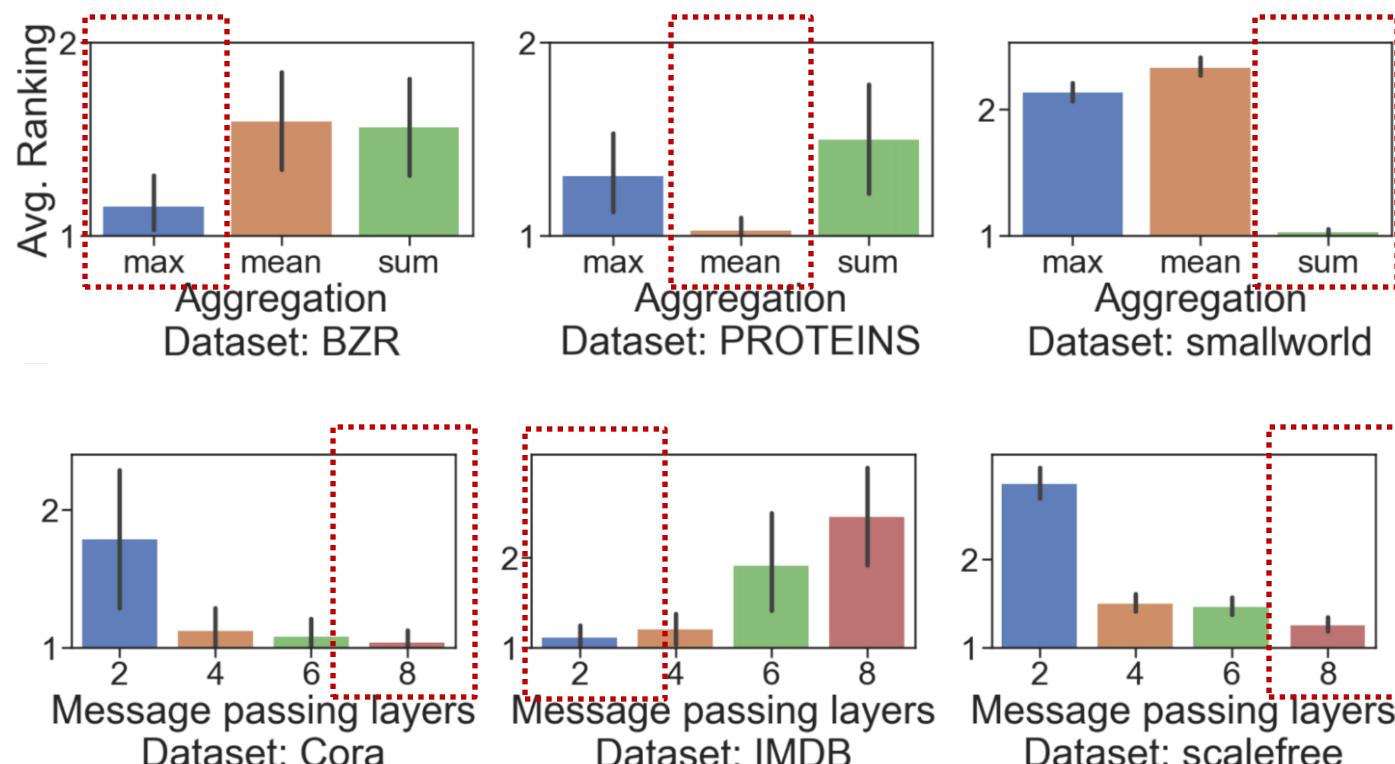
Highly dependent on the task



**Explanation:**  
Adam is more robust  
More training epochs is better

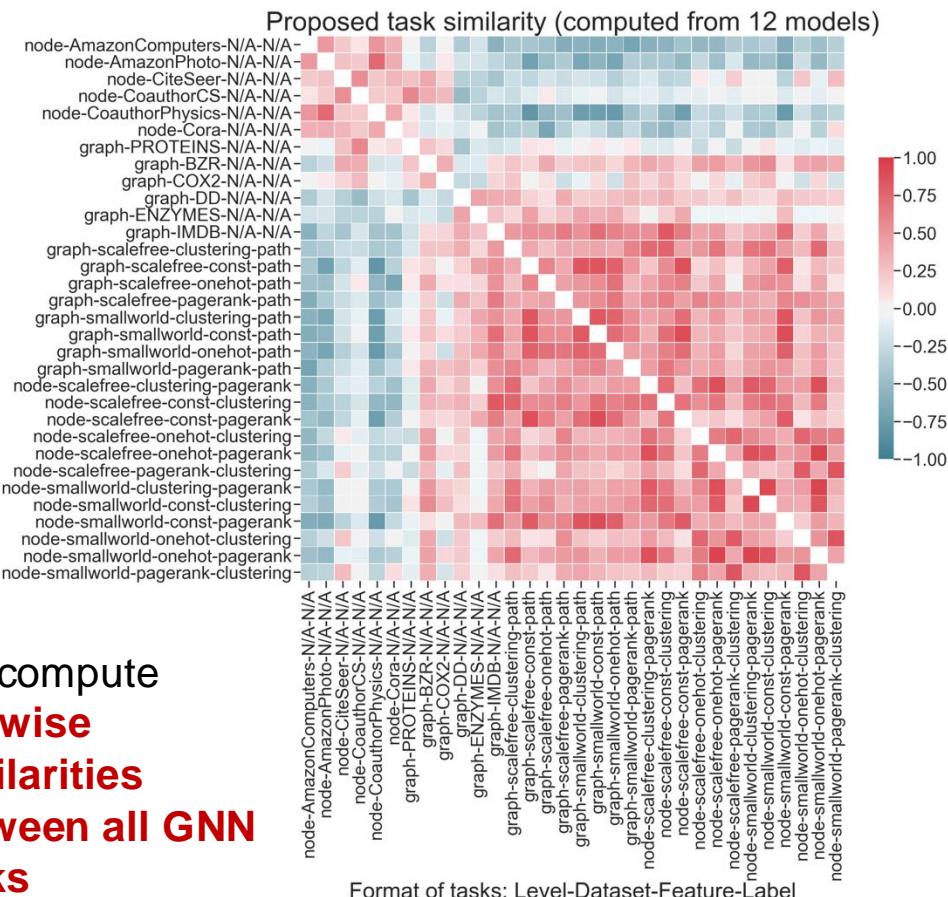
# Results 2: Understanding GNN Tasks

- Best GNN designs in different tasks vary significantly
  - Motivate that studying a task space is crucial



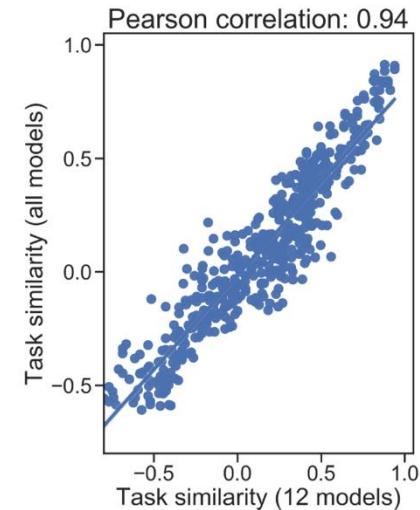
# Results 2: Understanding GNN Tasks

## Build a GNN task space



Recall how we **compute task similarity**

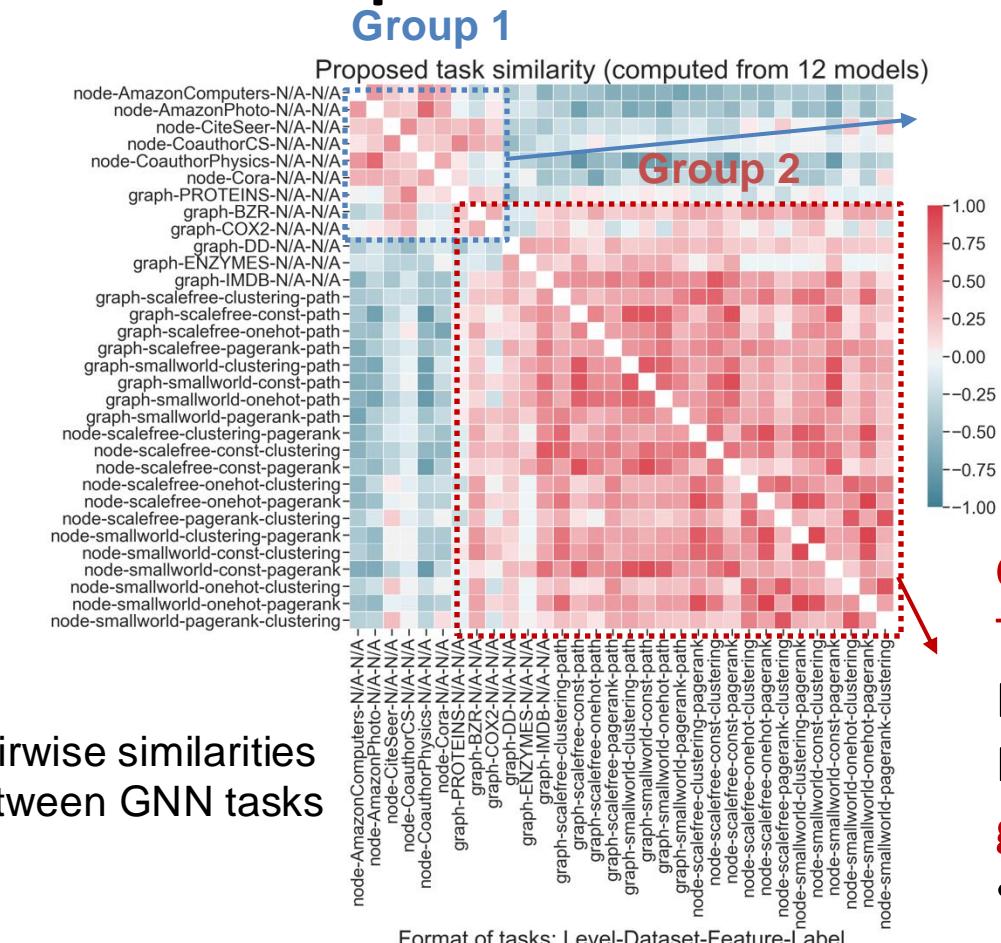
	Anchor Model Performance ranking					Similarity to Task A	
	Task A	$M_1$	$M_2$	$M_3$	$M_4$	$M_5$	1.0
Task B	Task B	$M_1$	$M_3$	$M_2$	$M_4$	$M_5$	0.8
Task C	Task C	$M_5$	$M_1$	$M_4$	$M_3$	$M_2$	-0.4



**Task similarity computation is cheap:  
Using 12 anchor models is a good approximation!**

# Results 2: Understanding GNN Tasks

- Proposed GNN task space is informative



Group 1:

Tasks rely on **feature information**  
Node/graph classification tasks, where  
**input graphs have high-dimensional  
features**

- Cora graph has 1000+ dim node feature

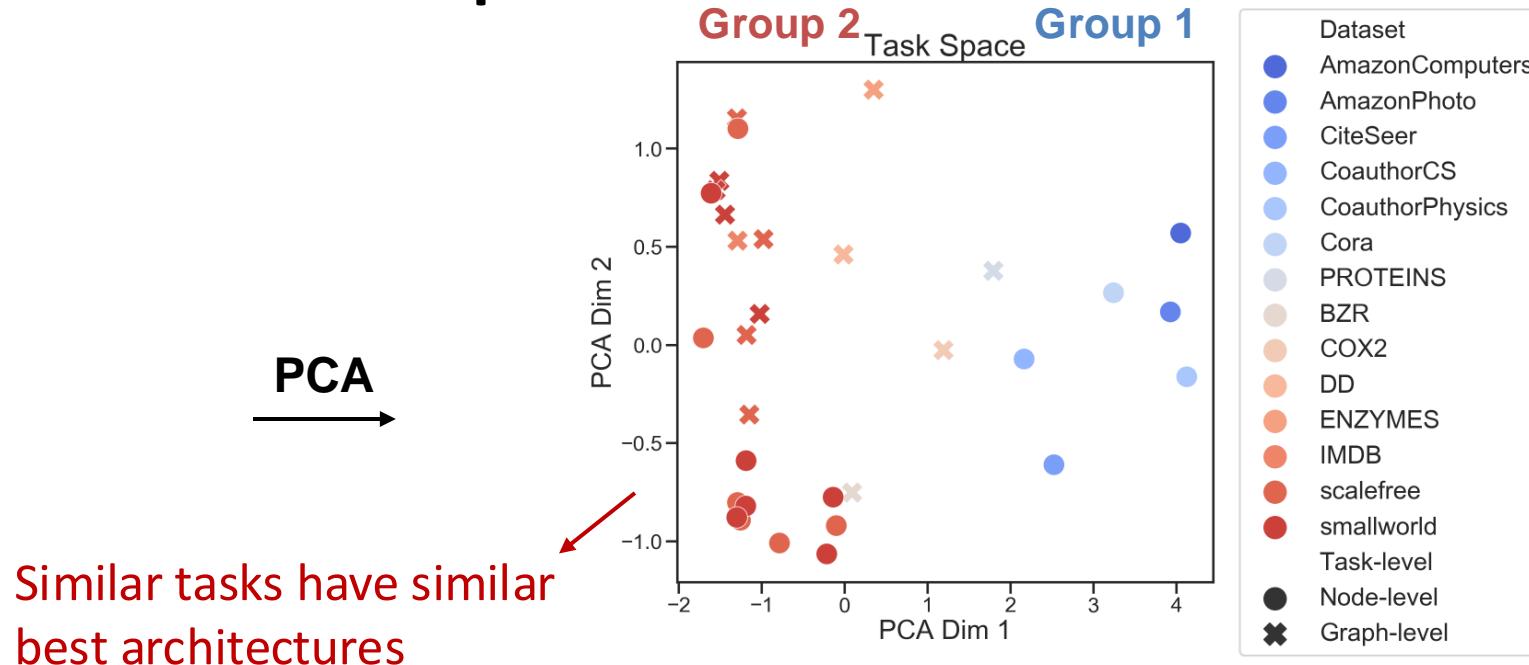
Group 2:

Tasks rely on **structural information**  
Nodes have few features  
Predictions are highly **dependent on  
graph structure**

- Predicting clustering coefficients

# Results 2: Understanding GNN Tasks

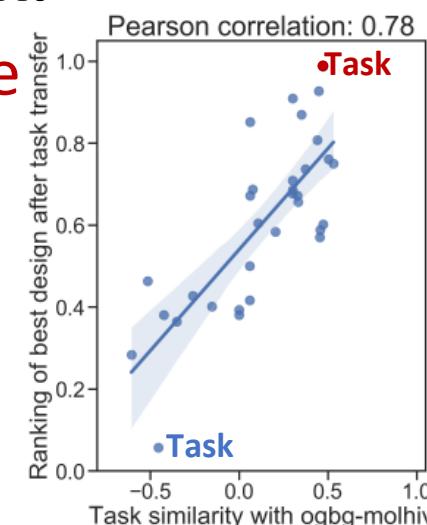
- Proposed GNN task space is informative



	Pre layers	MP layers	Post layers	Connectivity	AGG
Task A	2	8	2	skip-sum	sum
Task B	1	8	2	skip-sum	sum
Task C	2	6	2	skip-cat	mean

# Results 3: Transfer to Novel Tasks

- **Case study:** generalize best models to **unseen** OGB ogbg-molhiv task
  - **ogbg-molhiv is unique from other tasks:** 20x larger, imbalanced (1.4% positive) and requires out-of-distribution generalization
- **Concrete steps for applying to a novel task:**
  - **Step 1:** Measure 12 anchor model performance on the new task
  - **Step 2:** Compute similarity between the new task and existing tasks
  - **Step 3:** Recommend the best designs from existing tasks with high similarity



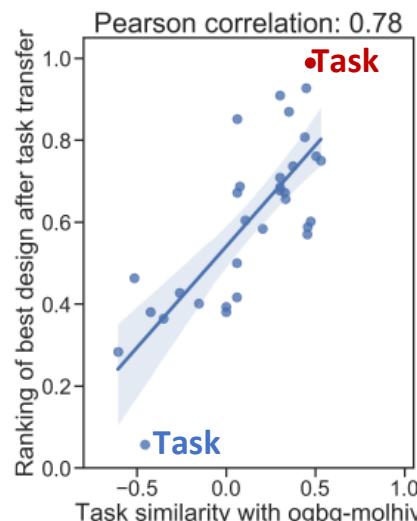
# Results 3: Transfer to Novel Tasks

- Our task space can guide best model transfer to novel tasks!

We pick 2 tasks:

Task A: Similar to OGB

Task B: Not similar to OGB



## Findings:

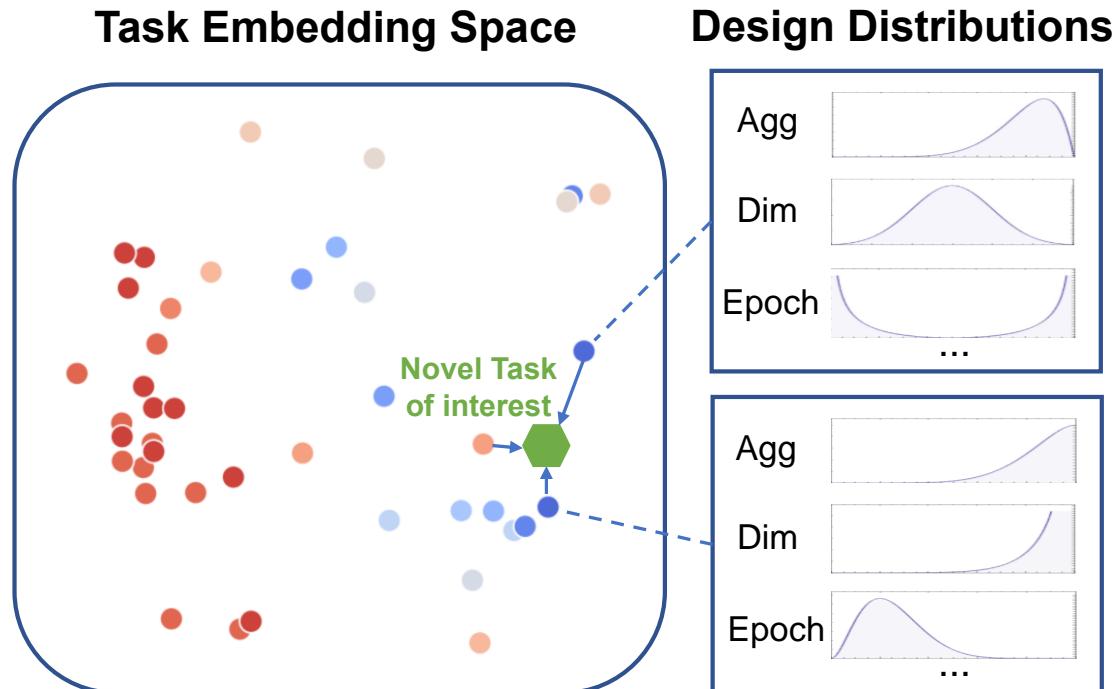
Transfer the best model from Task A achieves SOTA on OGB

Transfer the best model from Task B performs badly on OGB

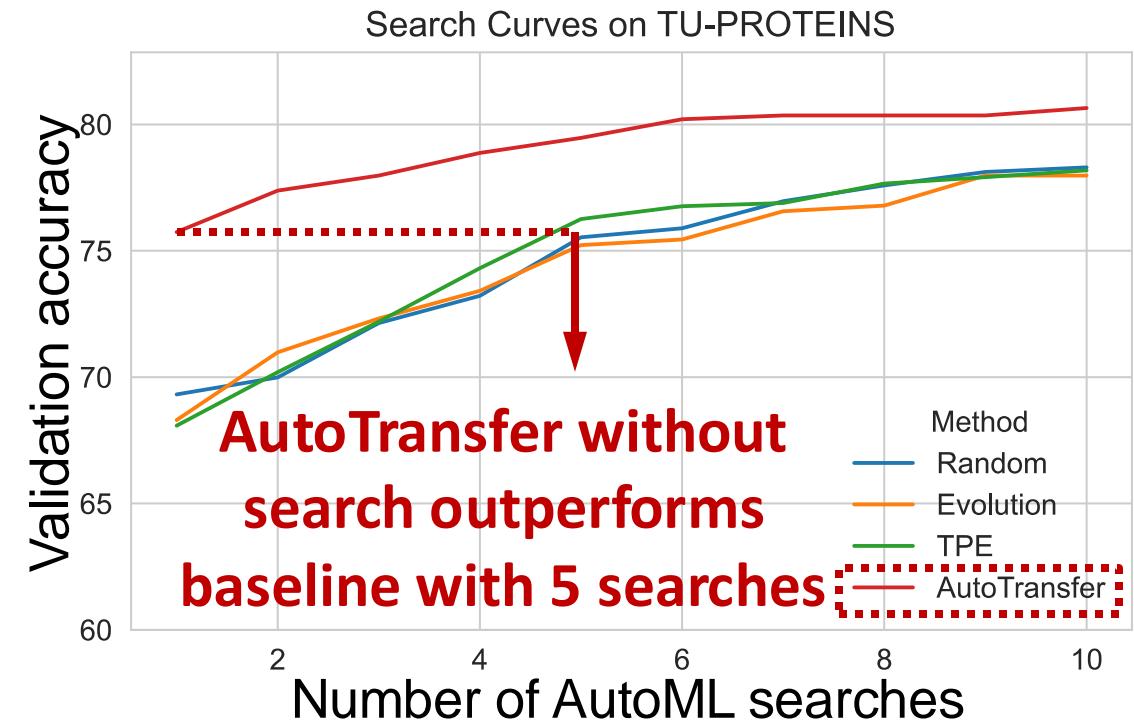
	Task A: graph-scalefree-const-path	Task B: node-CoauthorPhysics
Best design in our design space	(1, 8, 3, skipcat, sum)	(1, 4, 2, skipcat, max)
Task Similarity with ogbg-molhiv	0.47	-0.61
Performance after transfer to ogbg-molhiv	0.785	0.736

Previous SOTA: 0.771

# Extend to General AutoML Problems



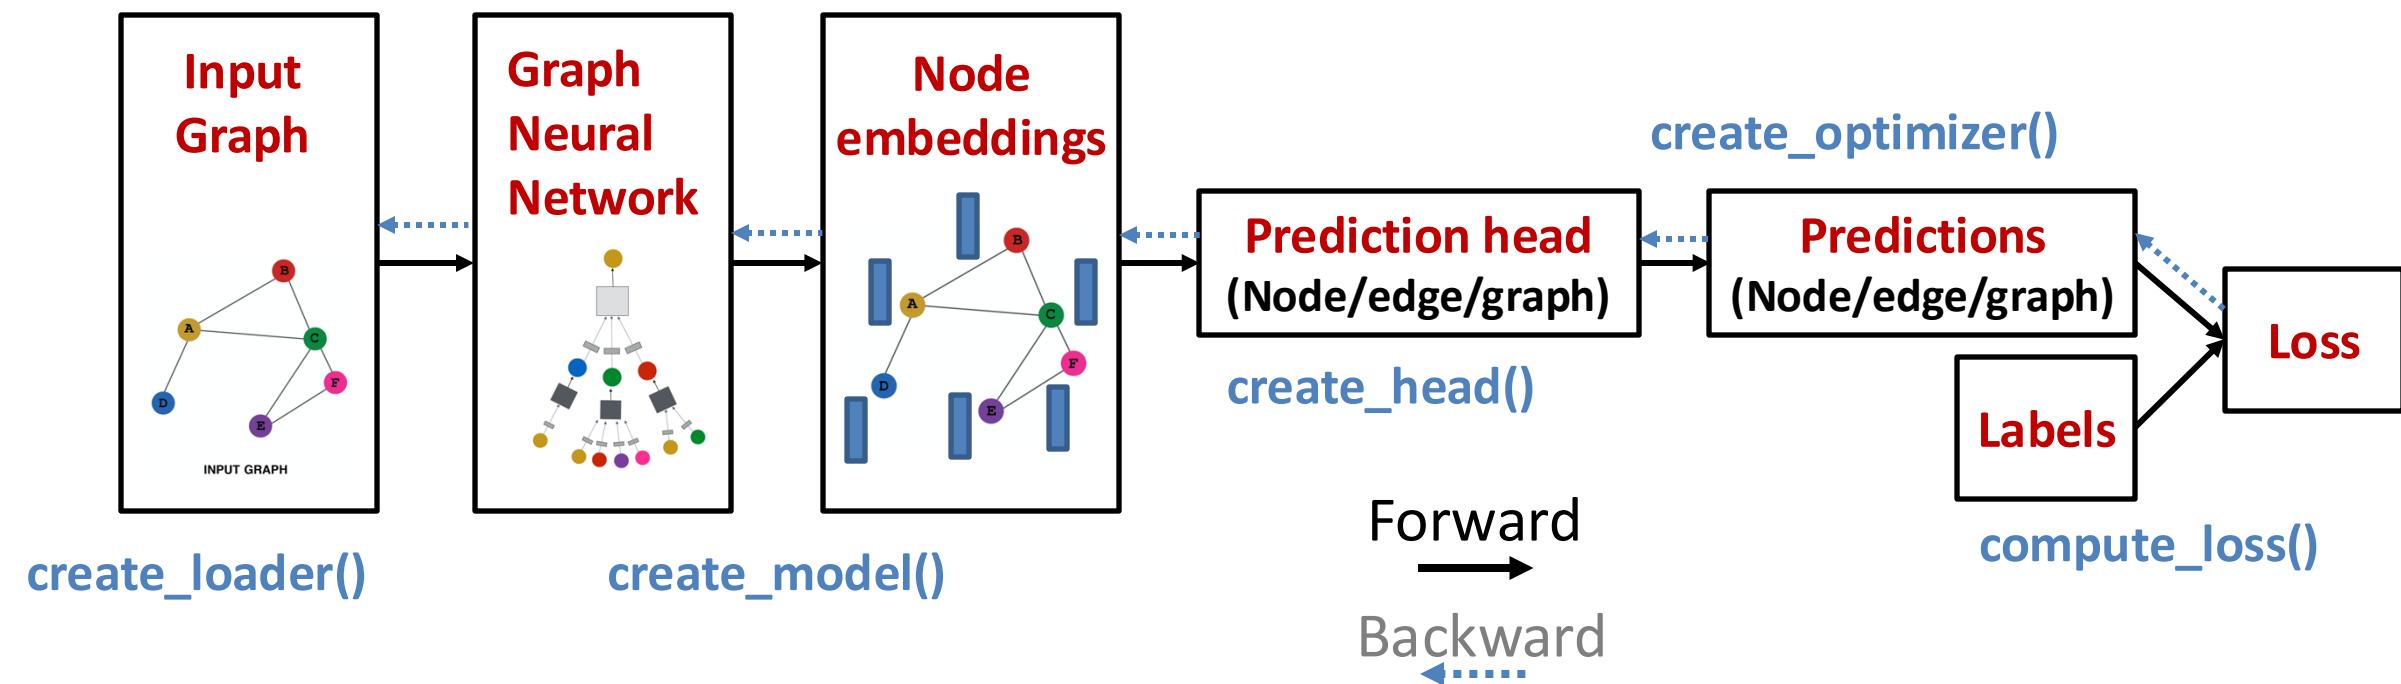
Sample from the design distributions of the most **similar tasks** through the **quantitative task similarity metric**



**Given a new task, AutoTransfer significantly saves the AutoML search cost**

# Software for Graph ML- GraphGym

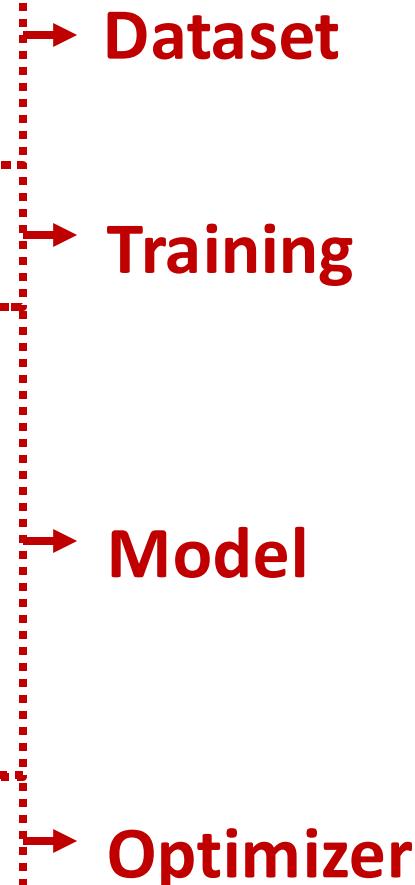
- **GraphGym:** powerful & easy-to-use primitives for graph ML
  - Users can quickly find the best GNN design for their specific tasks



# GraphGym: Easy-to-use API

## GraphGym config

```
1  out_dir: results
2  dataset:
3      format: PyG
4      name: Cora
5      task: node
6      task_type: classification
7      node_encoder: False
8      node_encoder_name: Atom
9      edge_encoder: False
10     edge_encoder_name: Bond
11
12 train:
13     batch_size: 128
14     eval_period: 1
15     ckpt_period: 100
16     sampler: full_batch
17
18 model:
19     type: gnn
20     loss_fun: cross_entropy
21     edge_decoding: dot
22     graph_pooling: add
23
24 gnn:
25     layers_pre_mp: 0
26     layers_mp: 2
27     layers_post_mp: 1
28     dim_inner: 16
29     layer_type: gcnconv
30     stage_type: stack
31     batchnorm: False
32     act: prelu
33     dropout: 0.1
34     agg: mean
35     normalize_adj: False
36
37 optim:
38     optimizer: adam
39     base_lr: 0.01
40     max_epoch: 200
```



- **GraphGym: A config fully describes an end-to-end GNN experiment**

- **Run an experiment in 1 line**

```
python main.py --cfg example_node.yaml --repeat 3
```

- **Run batch of experiments in 2 lines**

```
# generate configs
python configs_gen.py --config example.yaml --grid example.txt
# launch batch of experiments
bash run_batch.sh configs/example_grid_example 3 8
```

- **Democratize Graph ML for domain experts and practitioners**

# GraphGym Accelerates GNN Research

Proposed a **theoretically expressive graph ML model: ID-GNN**

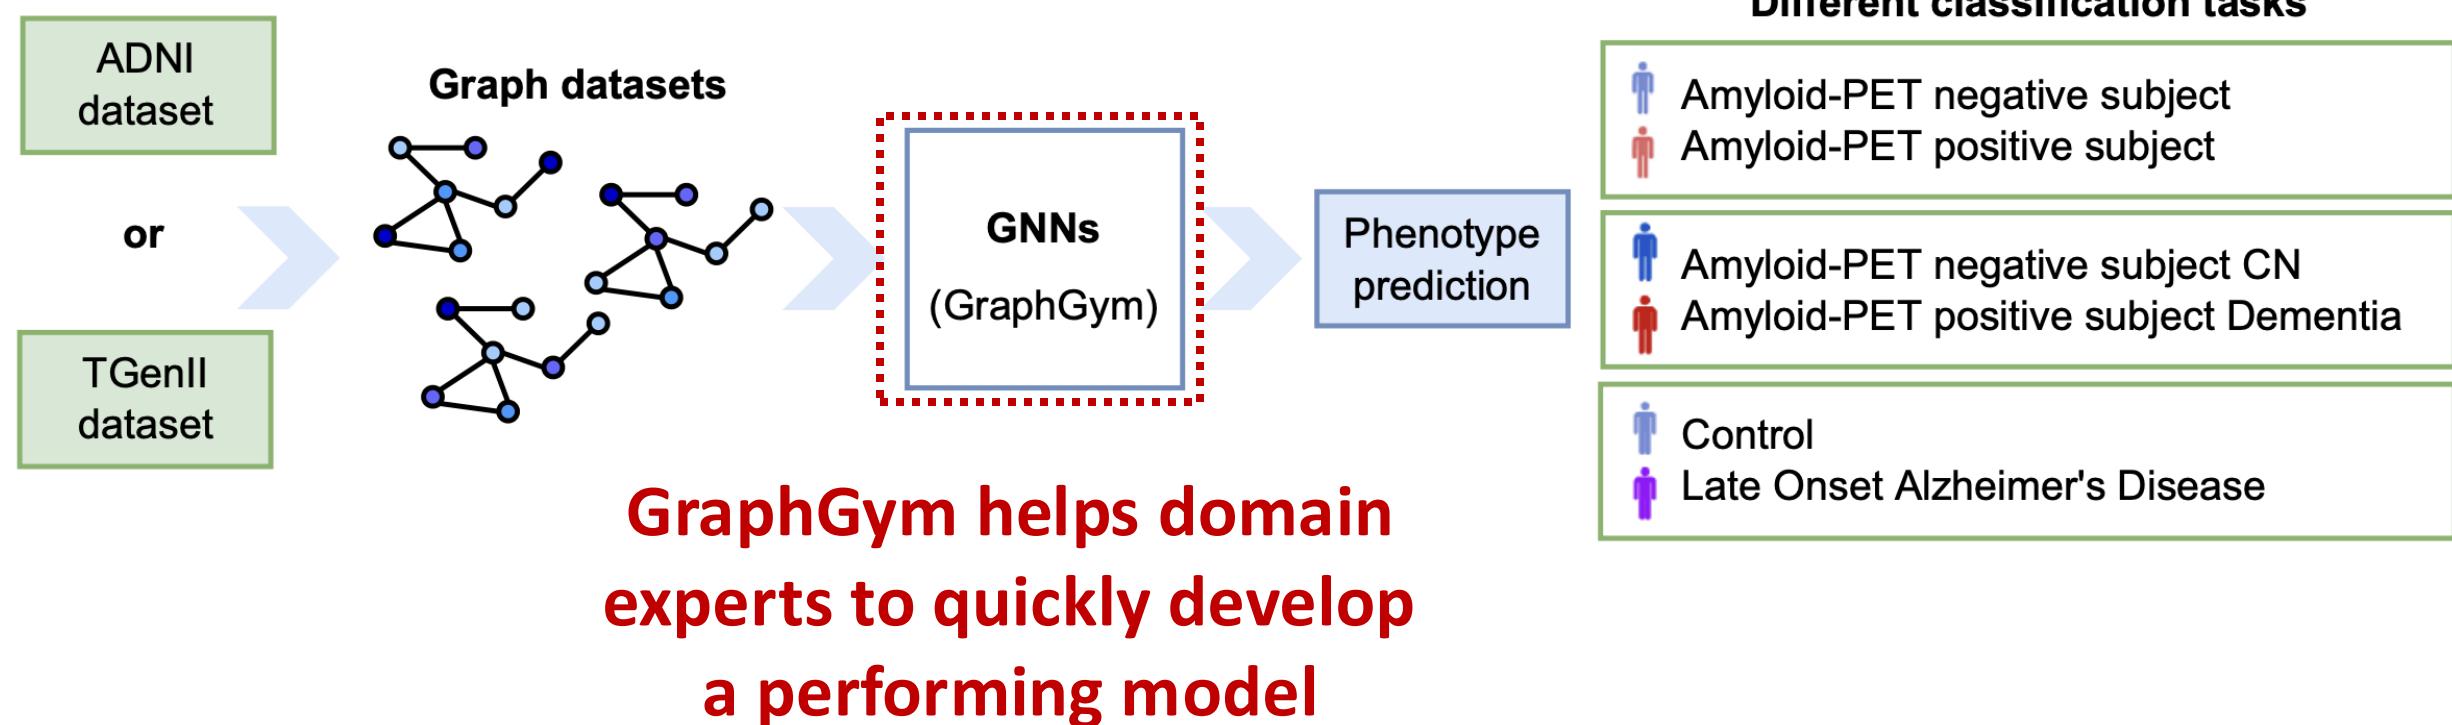
Experimentally show **ID-GNN** is better than **GNN**



- **2 lines of code with GraphGym**
  - **Accelerate research and promote knowledge sharing (reproducibility)**

```
# generate configs
python configs_gen.py --config baseline.yaml --config_budget baseline.yaml --grid idgnn.txt
# run batch of configs
# Args: config_dir, num of repeats, max jobs running, sleep time
bash run_batch.sh configs/baseline_grid_idgnn 3 10 1
```

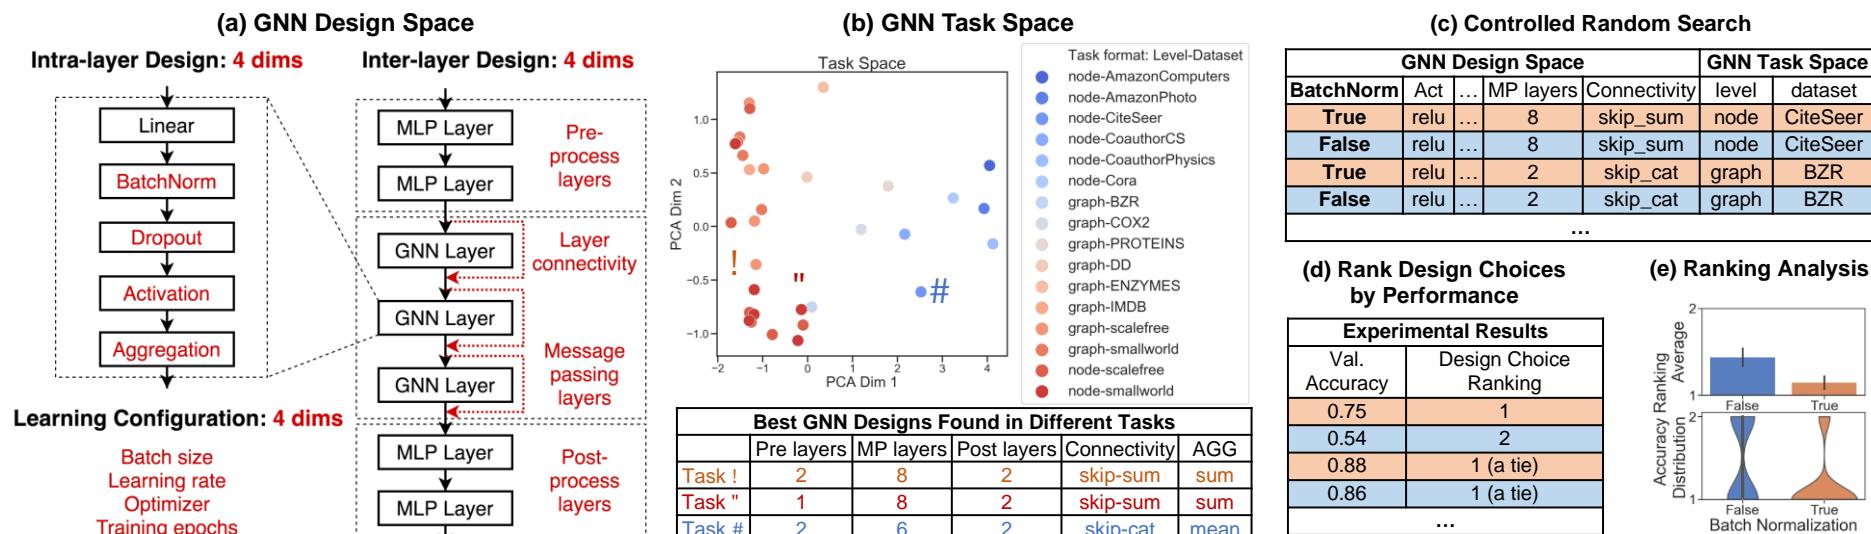
# GraphGym Accelerates Scientific Discovery

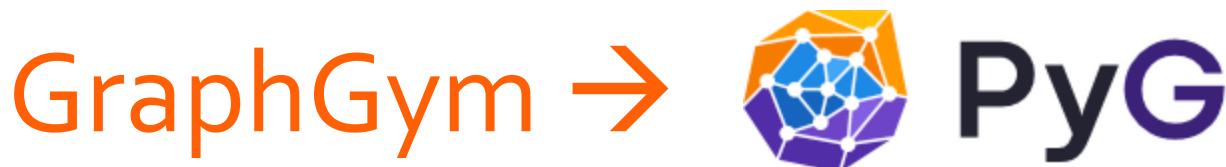


Hernández-Lorenzo, et al. "Graph Neural Networks for the Early Diagnosis of Alzheimer's Disease." *Nature Scientific Reports*, 2022

# GNN Design Space: Summary

- The first systematic investigation of:
  - General guidelines for GNN design
  - Understandings of GNN tasks
  - Transferring best GNN designs across tasks
  - **GraphGym:** Easy-to-use **code platform for GNN**





- GraphGym is integrated to PyG
- PyG: now **the most popular graph learning library**
  - **21,000+ Github Stars, 130,000+ monthly downloads**
  - The **go-to library for GNN researchers**: used in **1600 research papers**
  - **Widely used in industry**: Airbus, Amazon, Aramco, AstraZeneca, ...

```
from torch_geometric.nn import GCNConv

class GNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = GCNConv(3, 64)
        self.conv2 = GCNConv(64, 64)

    def forward(self, input, edge_index):
        h = self.conv1(input, edge_index)
        h = h.relu()
        h = self.conv2(h, edge_index)
        return h
```

GNN Implementation

PyTorch Geometric: Practical Realization of  
Graph Neural Networks

# Implementing Graph Neural Networks is challenging

- Sparsity *and* irregularity of the underlying data
  - How can we effectively parallelize irregular data of potentially varying size?
- Heterogeneity of the underlying data
  - numerical, categorical, image and text features, potentially over *different* types of data
- Inherently dynamic
  - it is hard to find scenarios in which graphs will *not* change over time
- Various *different* requests on scalability
  - sparse *vs.* dense graphs, many small *vs.* single giant graphs, ...
- Applicability to a set of *diverse* tasks
  - node-level *vs.* link-level *vs.* graph-level, clustering, pre-training, self-supervision, ...

# PyTorch Geometric

- **PyG (PyTorch Geometric)** is the PyTorch library to unify deep learning on graph-structured data.
  - simplifies implementing and working with Graph Neural Networks
  - bundles state-of-the-art GNN architectures and training procedures
  - achieves high GPU throughput on **sparse data** of varying size
  - suited for both academia and industry
    - flexible, comprehensive, easy-to-use

# Design Principles

- **Graph-based Neural Network Building Blocks**
  - Message Passing layers
  - Normalization layers
  - Pooling & Readout layers
- **Graph Transformations & Augmentations**
  - Graph diffusion
  - Missing feature value imputation
  - Mesh and Point Cloud support

# Design Principles

- **In-Memory Graph Storage, Datasets & Loaders**
  - Support for heterogeneous graphs 200+ benchmark datasets
  - 10+ sampling techniques
- **Examples & Tutorials**
  - Learn practically about GNNs
  - Videos, Colabs & Blogs
  - Application-driven Graph ML Tutorials
  - Notebooks examples with visualization

# Design Principles

- PyG is highly modular
  - Access to 200+ datasets and 50+ transforms
  - Access to a variety of mini-batch loaders
    - Node-wise sampling, subgraph-wise sampling, graph-wise batching
  - Access to 80+ GNN layers, normalizations and readouts as neural network building blocks 20+ pre-defined models
    - SAGEConv, GCNConv, GATConv, GINConv, PNACConv, ...
    - GraphSAGE, GCN, GAT, GIN, PNA, SchNet, DimeNet, ...
  - Access to regular PyTorch loss functions and training routines
    - Classification, Regression, Self-Supervision, ... Node-level, Link-level, Graph-level

# Design Principles

```
dataset = Reddit(root_dir, transform)

loader = NeighborLoader(dataset, num_neighbors=[25, 10])

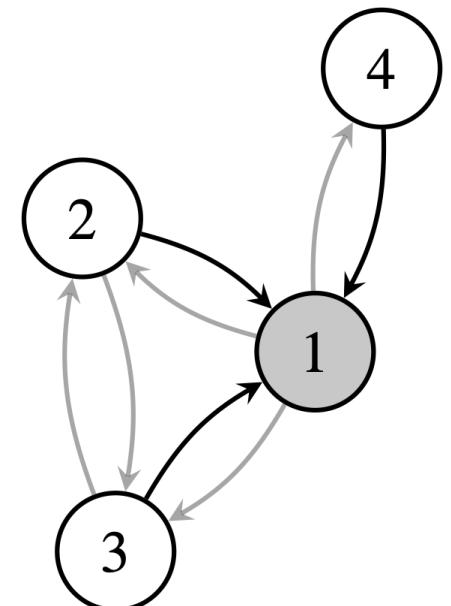
class GNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = SAGEConv(F_in, F_hidden)
        self.conv2 = SAGEConv(F_hidden, F_out)

    def forward(x, edge_index):
        x = self.conv1(x, edge_index)
        x = x.relu()
        x = self.conv2(x, edge_index)
        return x

    for data in loader:
        data = data.to(device)
        out = model(data.x, data.edge_index)
        loss = criterion(out, data.y)
        loss.backward()
        optimizer.step()
```

# Implementing Graph Neural Networks

- Given a sparse graph  $\mathcal{G} = \left( \mathbf{H}^{(0)}, (\mathbf{I}, \mathbf{E}) \right)$  with
  - Input node features  $\mathbf{H}^{(0)} \in \mathbb{R}^{|\mathcal{V}| \times C}$
  - Edge indices  $\mathbf{I} \in \{1, 2, \dots, |\mathcal{V}|\}^{2 \times |\mathcal{E}|}$
  - Optional edge features  $\mathbf{E} \in \mathbb{R}^{|\mathcal{E}| \times D}$
- Message Passing Scheme**
  - permutation-invariant aggregation operator,  
e.g., sum, mean or max
  - $\mathbf{h}_i^{(\ell+1)} = \text{Update}_{\theta}(\mathbf{h}_i^{(\ell)}, \bigoplus_{j \in \mathcal{N}(i)} \text{Message}_{\theta}(\mathbf{h}_i^{(\ell)}, \mathbf{h}_j^{(\ell)}, \mathbf{e}_{i,j}))$



# Graph Creation

- PyG expects a graph in either COO or CSR format to model complex (heterogeneous graphs)
- Both nodes and edges can hold any set of curated features
- **Benchmarking graphs:** Training labels and dataset splits are pre-defined
- **Real-world graphs in applications:** Build/Bring your own training labels and dataset splits



```
data = HeteroData()
data['user'].x = ...
data['product'].x = ...
data['user', 'product'].edge_index = ...
data['user', 'product'].edge_attr = ...
data['user', 'product'].time = ...
```

Best practices for  
graph design  
suitable for message  
passing

Best practices for  
feature selection

Best practices for  
feature encoding  
suitable for neural  
networks

# Task Formulation

- PyG supports any graph-related machine learning task, but curation of training labels is a user task



```
# Add user-level node labels:  
data['user'].y = ...  
  
transform = RandomNodeSplit(num_val=0.1, num_test=0.1)  
data = transform(data)
```

- **Benchmarking graphs:** Training labels and dataset splits are pre-defined
- **Real-world graphs in applications:** Build/Bring your own training labels and dataset splits

Best practices for curating training labels to map to ML task

Best practices for avoiding data leakage in temporal scenarios

Best practices for dataset splitting to match real-world use-cases

GNN Implementation

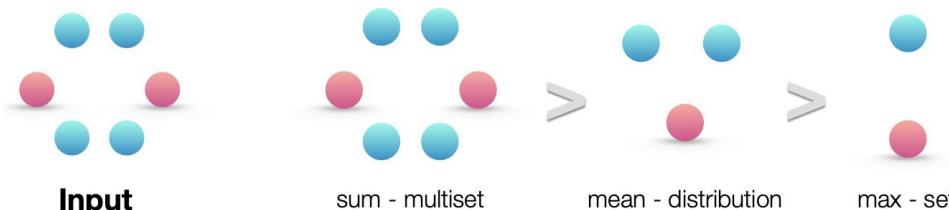
PyTorch Geometric: Advanced Features

# PyG Highlights

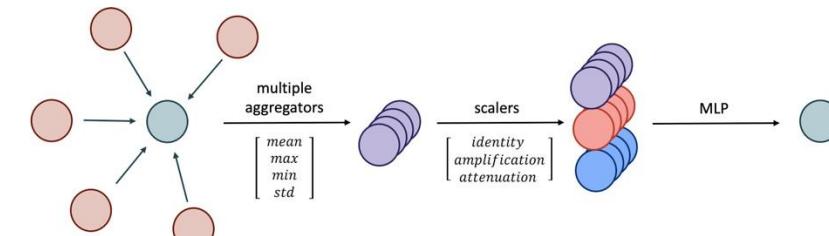
- Advanced Features
  - Improved GNN design via principled aggregations
  - Improved efficiency on heterogeneous graphs
  - Out-of-core data interfaces
  - Explainability and Compilation
  - GNN + LLM

# Principled Aggregations

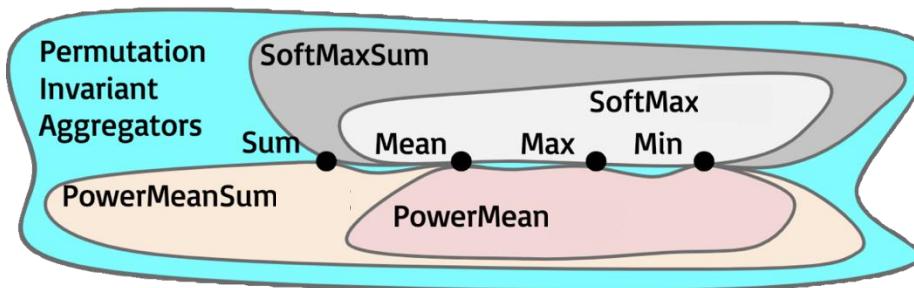
- Choice of neighborhood aggregation is a central topic in Graph ML research



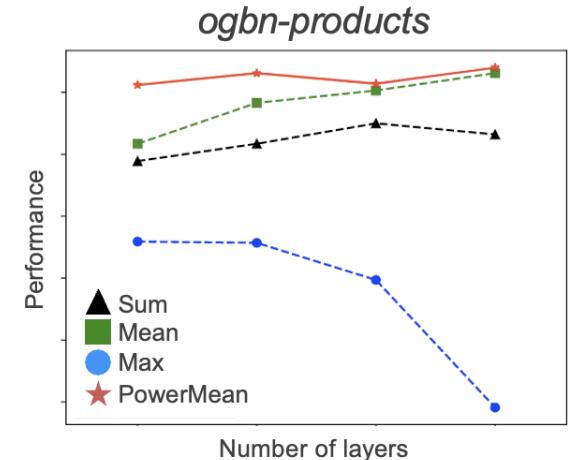
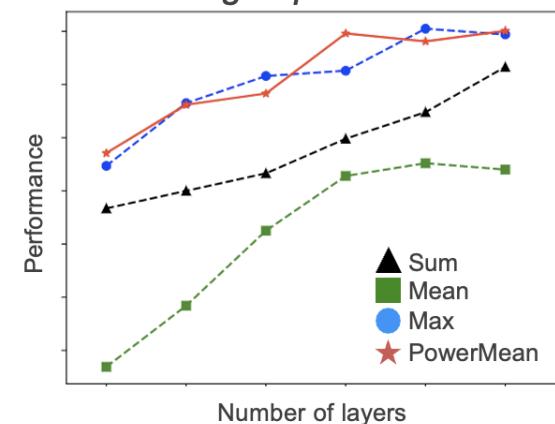
Xu et al.: How Powerful Are Graph Neural Networks?



Corso et al.: Principal Neighborhood Aggregation for Graph Nets

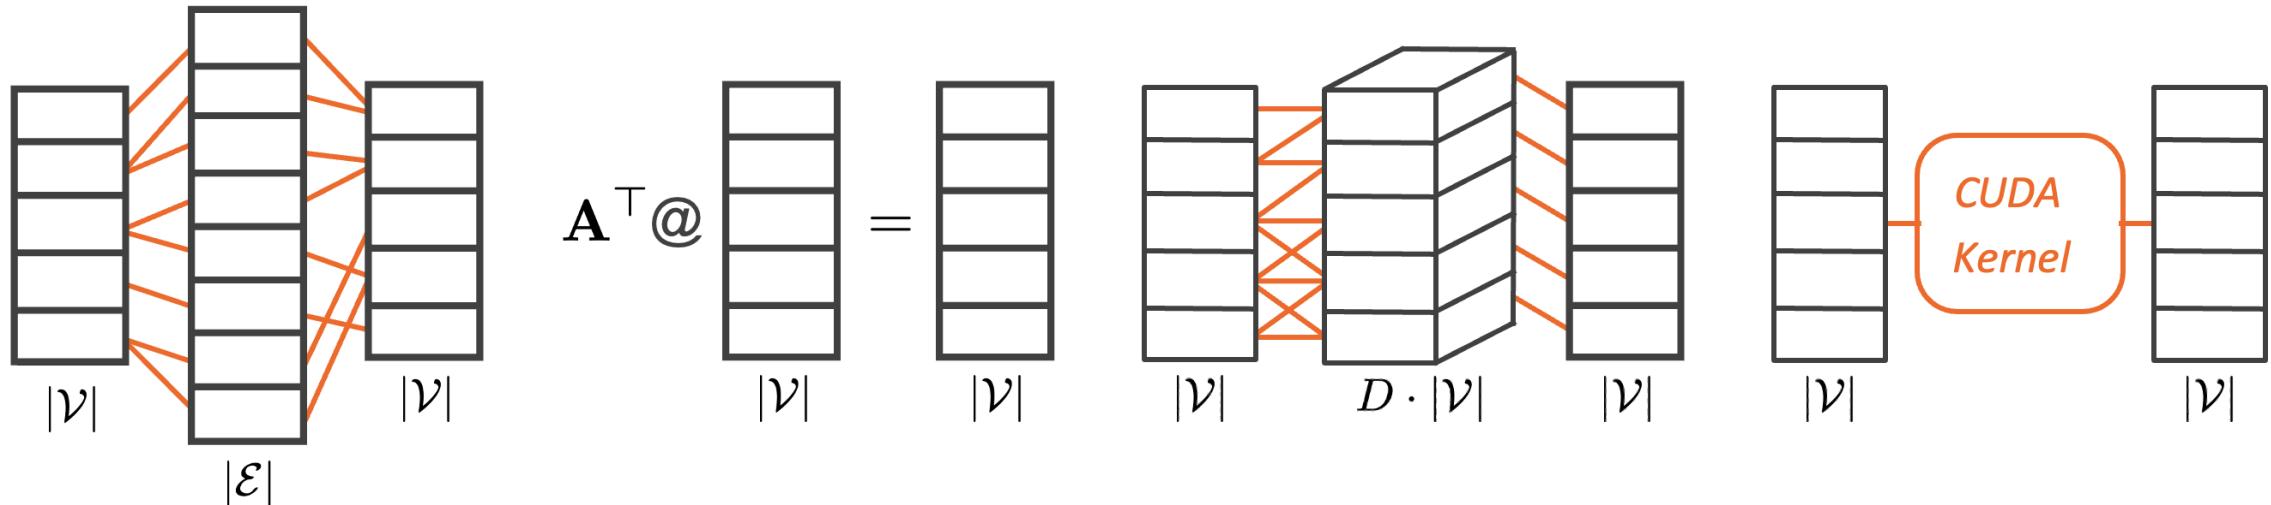


Li et al.: Deeper-GCN: All You Need to Train Deeper GCNs



# Principled Aggregations

The *different* flavors of implementing aggregations



## Gather & Scatter ( $\text{Py} \geq 0.1$ )

- very flexible 😊😊
- fast for sparse graphs 😊😊
- memory-inefficient 🤦‍♂️🤦‍♂️

## Sparse MatMul ( $\text{Py} \geq 1.6$ )

- less flexible 🤦‍♂️🤦‍♂️
- very fast 😊😊
- memory-efficient 😊😊

## Degree Bucketing ( $\text{Py} \geq 2.1$ )

- very flexible 😊😊
- fast for sparse graphs 😊😊
- memory-inefficient 🤦‍♂️🤦‍♂️

## Individual Kernel ( $\text{Py} \geq 2.2$ )

- not flexible at all 🤦‍♂️🤦‍♂️
- memory-efficient 😊😊
- very fast 😊😊

# Principled Aggregations

- PyG makes the concept of aggregations a first-class principle
  - Access to all kinds of simple, advanced, learnable and exotic aggregations
    - Median, Softmax, Attention, LSTM, ...
  - Fully-customize and combine aggregations within *MessagePassing* or for global pooling
  - Aggregations will pick up the best format to accelerate computation
    - scatter reductions, degree bucketing, ...
  - Further optimization via fusion
    - minimize I/O from global memory

# Principled Aggregations

```
● ● ●

# Simple aggregations:
mean_aggr = aggr.MeanAggregation()
max_aggr = aggr.MaxAggregation()

# Advanced aggregations:
median_aggr = aggr.MedianAggregation()

# Learnable aggregations:
softmax_aggr = aggr.SoftmaxAggregation(learn=True)
powermean_aggr = aggr.PowerMeanAggregation(learn=True)

# Exotic aggregations:
lstm_aggr = aggr.LSTMAggregation()
sort_aggr = aggr.SortAggregation(k=4)

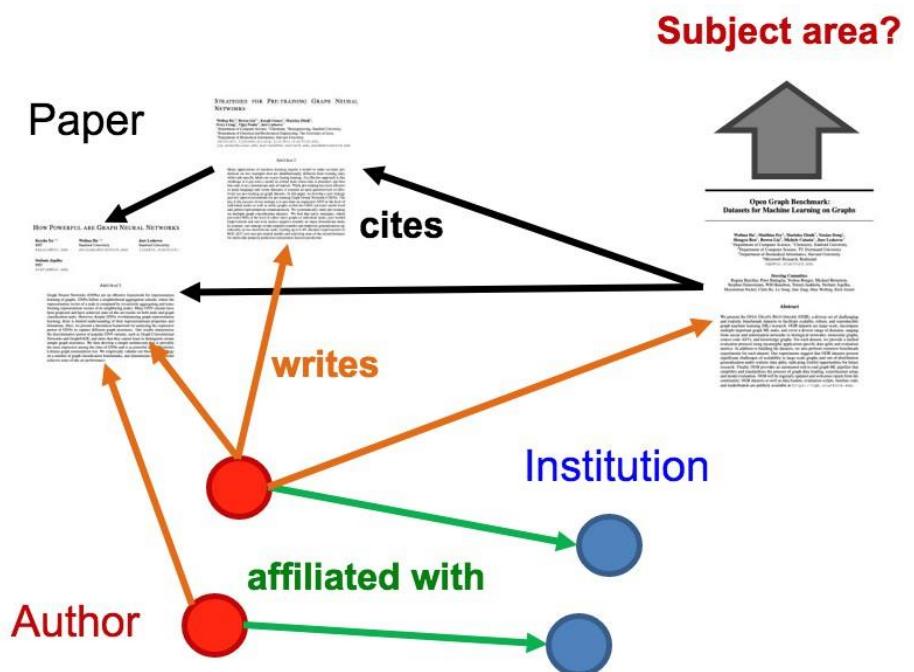
# Use within message passing:
conv = MyConv(aggr=[median_aggr, lstm_aggr])

# Use for global pooling:
h_graph = sort_aggr(h_node, batch)
```

# PyG Highlights

- Advanced Features
  - Improved GNN design via principled aggregations
  - **Improved efficiency on heterogeneous graphs**
  - Out-of-core data interfaces
  - Explainability and Compilation
  - GNN + LLM

# Heterogeneous Graph Support



- (Nearly) all real-world graphs are heterogeneous!

# Heterogeneous Graph Support

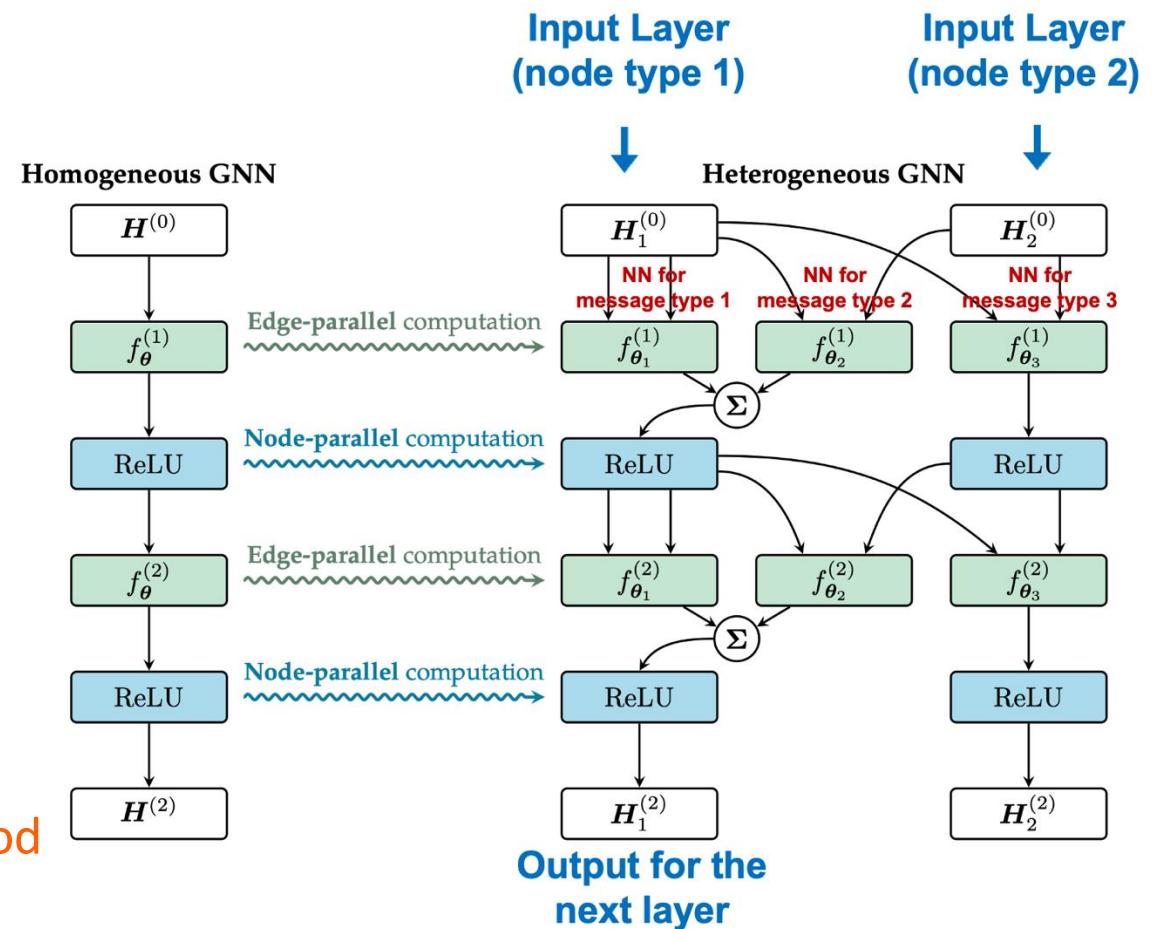
- Heterogeneous Graph Neural Networks

- A homogeneous GNN can be converted to a heterogeneous one by learning distinct parameters for each individual edge type

Edge type dependent parameters

$$\mathbf{h}_i^{(\ell+1)} = \sum_{r \in \mathcal{R}} f_{\theta_r}^{(\ell+1)}(\mathbf{h}_i^{(\ell)}, \{\mathbf{h}_j^{(\ell)} : j \in \mathcal{N}^{(r)}(i)\})$$

The number of relations  
Relational-wise neighborhood



# Heterogeneous Graph Support

- PyG can automatically convert homogeneous GNNs to heterogeneous ones
    - `to_hetero(model (node_types, edge_types))`: *Implemented via*
    - `to_hetero_with_bases(model (node_types, edge_types))`: *torch.fx*
- 1.Duplicates message passing modules for *each* edge type
  - 2.Transforms the underlying computation graph so that messages are exchanged along *different* edge types
  - 3.Uses lazy initialization (-1) to handle *different* input feature dimensionalities



```
from torch_geometric.nn import GAT, to_hetero

model = GAT(in_channels=-1, hidden_channels=64,
            out_channels=72, num_layers=2)

model = to_hetero(model, (node_types, edge_types))

out = model(data.x_dict, data.edge_index_dict)
```

# PyG Highlights

- Advanced Features
  - Improved GNN design via principled aggregations
  - Improved efficiency on heterogeneous graphs
  - **Out-of-core data interfaces**
  - Explainability and Compilation
  - GNN + LLM

# Out-of-core Data Interfaces

There exists multiple ways to scale PyG beyond single-node in-memory datasets

## 1. Pre-process subgraphs on disk and load them on-the-fly,

e.g., via dedicated Spark routines

- Naturally supported via PyG's mini-batch DataLoader
- No sampling/feature fetching overhead during GNN training
- Heavy pre-processing and storage requirements
- Experimentation with different sampling strategies and parameters gets increasingly harder

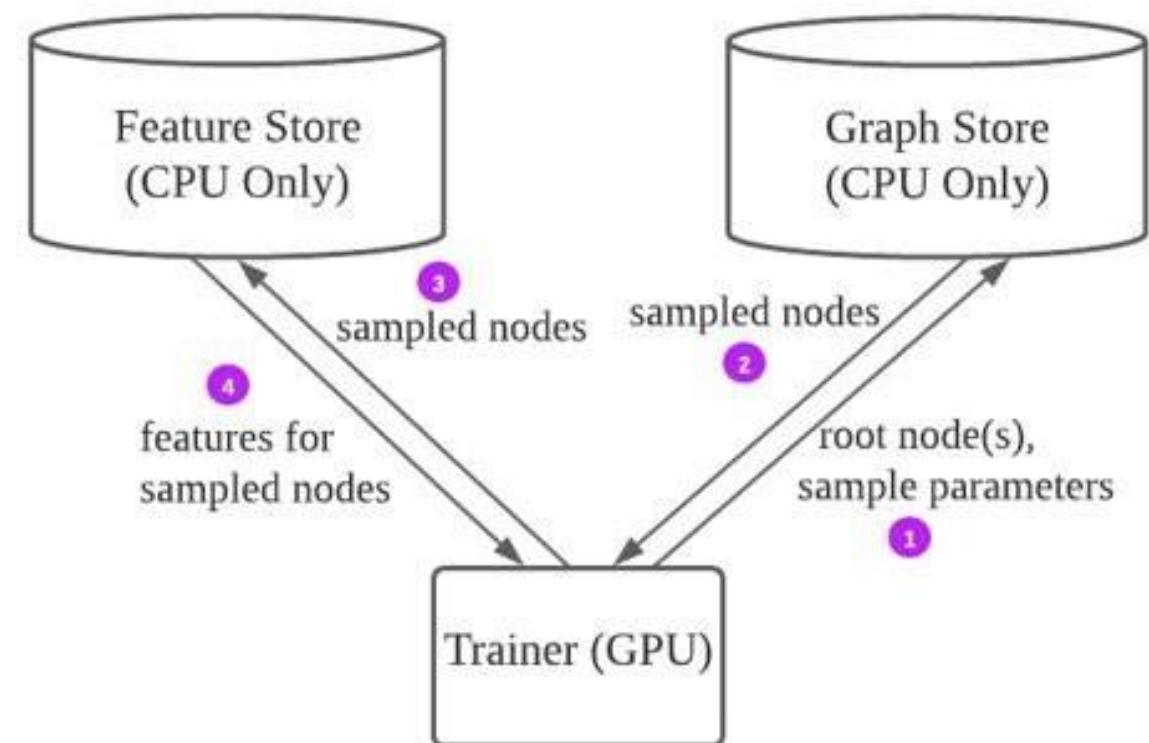
```
class Dataset:  
    def __getitem__(self, idx) -> Data:  
        # 1. Load subgraph from disk:  
        data = load()  
        # 2. Convert to PyG format:  
        data = Data()  
        return data  
  
dataset = Dataset()  
loader = DataLoader(dataset, batch_size=128)
```

# Out-of-core Data Interfaces

There exists multiple ways to scale PyG beyond single-node in-memory datasets

2. With PyG, we support any backend by providing FeatureStore and GraphStore abstractions

- Disentangles feature fetching from graph sampling routines
- Allows for distributed server/client architectures
- Allows for out-of-memory backends, e.g., via connecting to graph databases



# Out-of-core Data Interfaces

There exists multiple ways to scale PyG beyond single-node in-memory datasets

2. With PyG, we support any backend by providing FeatureStore and GraphStore abstractions

- These interfaces can enable an elastic, heterogeneous compute architecture
  - Read-optimized on-disk feature store via RocksDB
  - Low-latency in-memory graph store

*We are working with multiple graph database vendors to enable store implementations, e.g., KùzuDB*

```
● ● ●

class MyFeatureStore(FeatureStore):
    def get_tensor(self, attr):
        pass # Implement feature access

class MyGraphStore(GraphStore):
    def sample_from_nodes(self, index):
        pass # Implement node-wise sampling

    def sample_from_edges(self, index):
        pass # Implement edge-wise sampling
```

# PyG Highlights

- Advanced Features
  - Improved GNN design via principled aggregations
  - Improved efficiency on heterogeneous graphs
  - Out-of-core data interfaces
  - **Explainability and Compilation**
  - GNN + LLM

# Explainability

- PyG 2.3 introduced a new unified explainability interface to explain any GNN out-of-the-box across various explainer algorithms
  - Works on both homogeneous and heterogeneous graphs
  - Support for general explainers such as integrated gradients, saliency, shapley, etc
  - Support for dedicated GNN explainers, e.g., GNNExplainer, PGExplainer, etc
  - Support for a diverse range of tasks
  - Support for visualizations and metric computation



```
explainer = Explainer(  
    model=model,  
    algorithm=GNNExplainer(epochs=200),  
    node_mask_type='attributes',  
    edge_mask_type='object',  
    model_config=dict(  
        mode='multiclass_classification',  
        task_level='node',  
        return_type='probs',  
    ),  
)  
  
explanation = explainer(x, edge_index)
```

# Compilation

- PyG 2.3 takes full advantage of PyTorch 2.0 compilation, which makes your GNN run faster by JIT-compiling it into optimized kernels

Model	Mode	Forward	Backward	Total	Speedup	
GCN	Eager	2.6396s	2.1697s	4.8093s		 model = GraphSAGE() model = torch.compile(model)  out = model(data.x, data.edge_index)
GCN	Compiled	1.1082s	0.5896s	1.6978s	2.83x	
GraphSAGE	Eager	1.6023s	1.6428s	3.2451s		
GraphSAGE	Compiled	0.7033s	0.7465s	1.4498s	2.24x	
GIN	Eager	1.6701s	1.6990s	3.3690s		
GIN	Compiled	0.7320s	0.7407s	1.4727s	2.29x	

# PyG Highlights

- Advanced Features
  - Improved GNN design via principled aggregations
  - Improved efficiency on heterogeneous graphs
  - Out-of-core data interfaces
  - Explainability and Compilation
  - **GNN + LLM**

# PyG 2.6: Combining GNNs with LLMs

- In order to facilitate further research on combining GNNs with LLMs, **PyG 2.6** introduces:
  - A new sub-package `torch_geometric.nn.nlp` with fast access to SentenceTransformer models and LLMs
  - A new model `GRetriever` that is able to co-train LLAMA2 with GAT for answering questions based on knowledge graph information
  - A new example folder `examples/LLM` that shows how to utilize these models in practice

# Summary for Today

- **GraphGym:** Easy-to-use **code platform for GNN**
  - General guidelines for GNN design
  - Understandings of GNN tasks
  - Transferring best GNN designs across tasks
- **PyG principles & highlights**